

# *The Five Generations of Entity Resolution on Web Data*



ICWE 2024, Finland  
Monday, 17<sup>th</sup> of June

<https://edu.nl/97b8v>

**Konstantinos Nikoletos**

University of Athens  
Greece

[k.nikoletos@di.uoa.gr](mailto:k.nikoletos@di.uoa.gr)

**Ekaterini Ioannou**

University of Tilburg  
Netherlands

[ekaterinil.ioannou@uvt.nl](mailto:ekaterinil.ioannou@uvt.nl)

**George Papadakis**

University of Athens  
Greece

[gpapadis@di.uoa.gr](mailto:gpapadis@di.uoa.gr)

# Structure Outline

- Introduction
- Generations: 1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup>, 5<sup>th</sup>
- Hands-on Session
- Challenges and Final Remarks

# Part A – Introduction

- Motivation
- Preliminaries
- Generations: 1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup>, 5<sup>th</sup>
- Hands-on Session
- Challenges and Final Remarks

# Motivation

- Entities invaluable asset for numerous current applications and systems
- Encode a large part of our **knowledge**

## Matching, Linkage, Reconciliation, etc.

- Many names, descriptions, or IDs (URIs) are used for the same **real-world “entity”**




London 런던  
लंदन  
लंदन  
ロンドン  
লন্ডন  
லண்டன்  
இலண்டன்  
ਲੰਦਨ

## Disambiguation, Deduplication, etc.

- Plethora of different “entities” have the same name



London Island, Chile

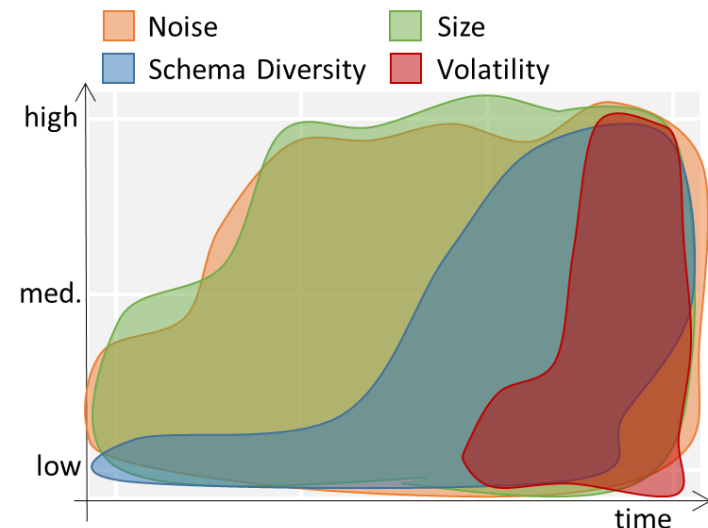



London, Jack  
2612 Almes Dr  
Montgomery, AL  
(334) 272-7005

# Motivation

**Entity Resolution** is required for data integration, link discovery, query answering, Web / object-oriented searching, etc.

- Goal remains the same for the last 50+ years
- BUT the challenges to be addressed are constantly evolving



cf. book: "The Four Generations of Entity Resolution"

# Entity Resolution

---

- Identifies and aggregates the **different** entity profiles that describe the **same** objects [1, 2, 3, 4]
- Primary usefulness:
  - Improves data quality and integrity
  - Fosters re-use of existing data sources
- Example application domains:
  - Linked Data
  - Building Knowledge Graphs
  - Census data
  - Price comparison portals



# Types of Entity Resolution

---

- The given entity collections can be of two types:  
**clean** + **dirty** [3,5]
- **Clean:**
  - Duplicate-free data
  - E.g., DBLP, ACM Digital Library, Wikipedia, Freebase
- **Dirty:**
  - Contain duplicate entity profiles
  - E.g., Google Scholar, CiteseerX

# Types of Entity Resolution


- The given entity collections can be of two types:  
**clean** + **dirty** [3,5]
- **Clean:**
  - Duplicate-free data
  - E.g., DBLP, ACM Digital Library, Wikipedia, Freebase
- **Dirty:**
  - Contain duplicate entity profiles
  - E.g., Google Scholar, CiteseerX





# Types of Entity Resolution

- Based on the quality of input, we distinguish entity resolution into 3 sub-tasks:
  - Clean-Clean ER (a.k.a. **Record Linkage** in databases)
  - Dirty-Clean ER
  - Dirty-Dirty EREquivalent to **Dirty ER** (a.k.a. **Deduplication** in databases)



# References

---


1. X. L. Dong, D. Srivastava. Big Data Integration. Synthesis Lectures on Data Management, Morgan & Claypool Publishers 2015, pp. 1-198.
2. A. K. Elmagarmid, P. G. Ipeirotis, V. S. Verykios. Duplicate Record Detection: A Survey. IEEE Trans. Knowl. Data Eng. 19(1): 1-16 (2007).
3. V. Christophides, V. Efthymiou, K. Stefanidis. Entity Resolution in the Web of Data. Synthesis Lectures on the Semantic Web: Theory and Technology, Morgan & Claypool Publishers 2015.
4. P. Christen. Data Matching - Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection. Data-Centric Systems and Applications, Springer 2012, ISBN 978-3-642-31163-5, pp. I-XIX, 1-270.
5. P. Christen. A Survey of Indexing Techniques for Scalable Record Linkage and Deduplication. IEEE Trans. Knowl. Data Eng. 24(9): 1537-1555 (2012).

- Introduction

## Part B – Generations

- Generation 1: tackling Veracity
  - Generation 2: tackling Volume and Veracity
  - Generation 3: tackling Variety, Volume and Veracity
  - Generation 4: tackling Velocity, Variety,  
Volume and Veracity
  - Generation 5: Entity Resolution Revisited:  
Leveraging External Knowledge
- Hands-on Session
  - Challenges and Final Remarks

# Generation 1: Tackling Veracity



- Earliest approach
- Scope:
  - Structured data
- Goal:
  - Achieve **high accuracy** despite inconsistencies, noise, or errors in entity profiles
- Assumptions:
  - Known schema → custom, schema-based solutions

# Step 1: Schema Alignment / Matching


- Scope:
  - Record Linkage
- Goal:
  - Create **mappings between equivalent attributes** of the two schemata, e.g., *profession*  $\equiv$  *job*
- Types of Solutions:
  - Structure-based
  - Instance-based
  - Usage-based
  - Hybrid

# Step 1: Schema Alignment / Matching

- Taxonomy of Main Schema Matching Methods (in chronological order)

Method	Category	Type of Evidence
Cupid [1]	Structure-based	Name similarity, Constraints, Contextual similarity
Similarity Flooding [2]	Structure-based	Name similarity, Contextual similarity
COMA [3]	Hybrid	Name similarity, Constraints, Contextual similarity
Distribution-based [4]	Instance-based	Value distribution

# Step 2: Blocking



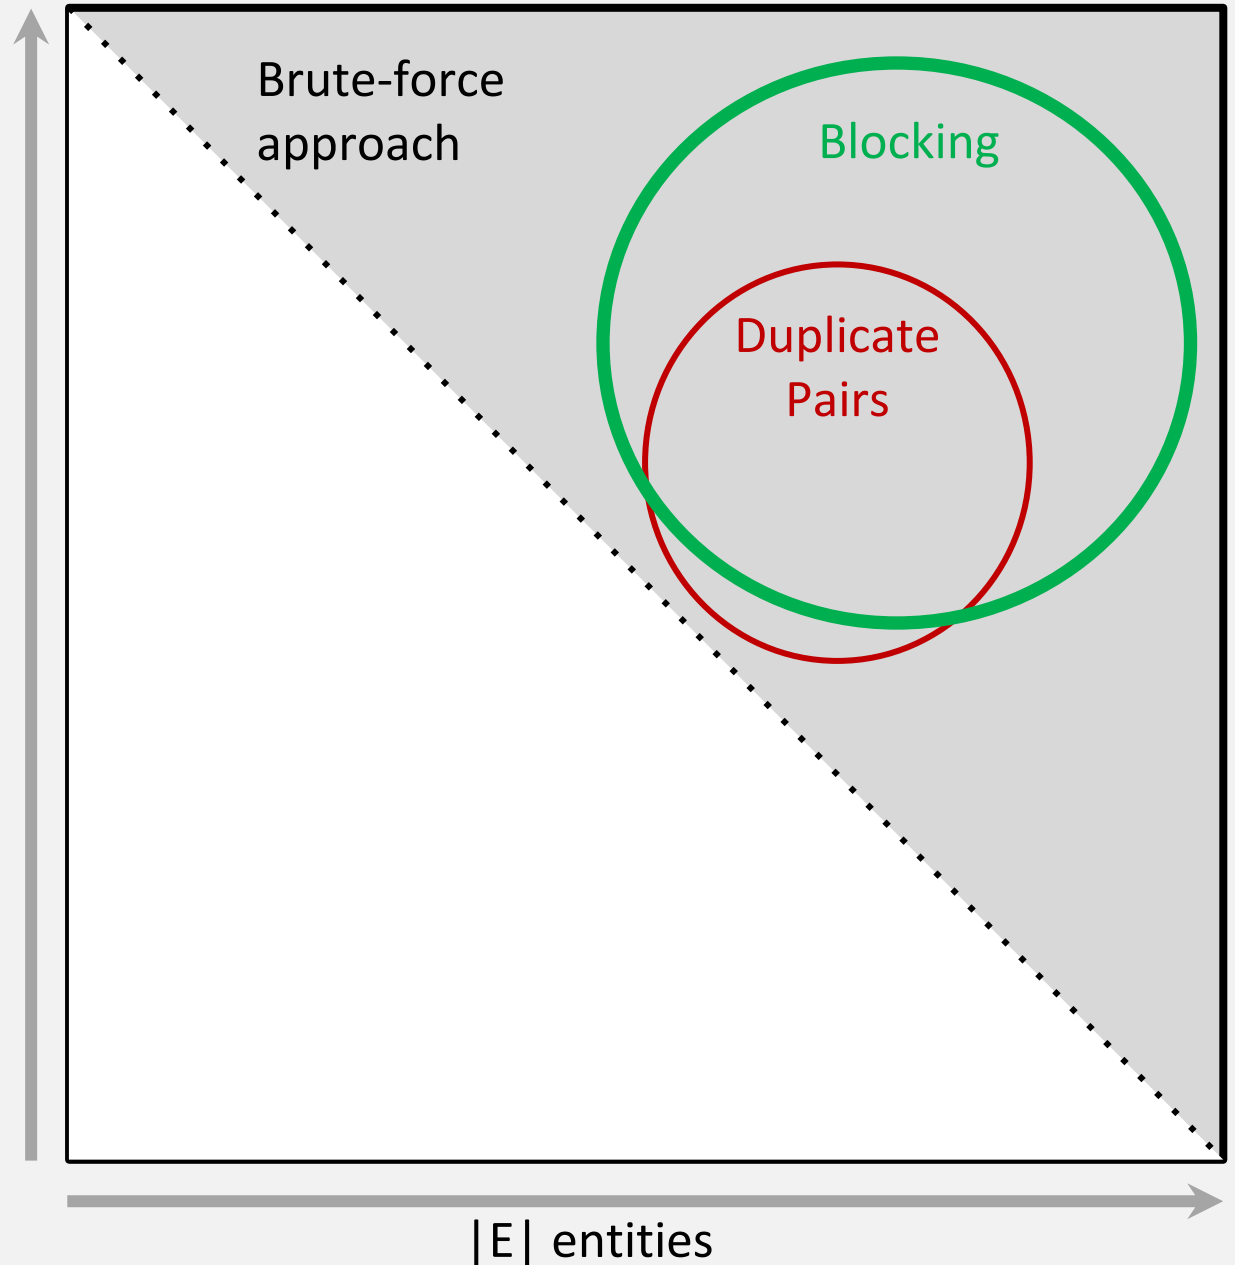
- Scope:
  - Both Deduplication and Record Linkage
- Goal:
  - ER is an inherently **quadratic** problem,  $O(n^2)$ : every entity has to be compared with all others
  - Blocking groups **similar** entities into **blocks**
    - Comparisons are executed **only inside** each block
    - Complexity is now quadratic to the size of the block (much smaller than dataset size!)

# Computational cost

Input:  
Entity Collection E

|E| entities

E.g.: For a dataset with  
**100,000** entities:  
 $\sim 10^{10}$  comparisons,  
If **0.05 msec** each  $\rightarrow$   
**>100 hours** in total



|E| entities



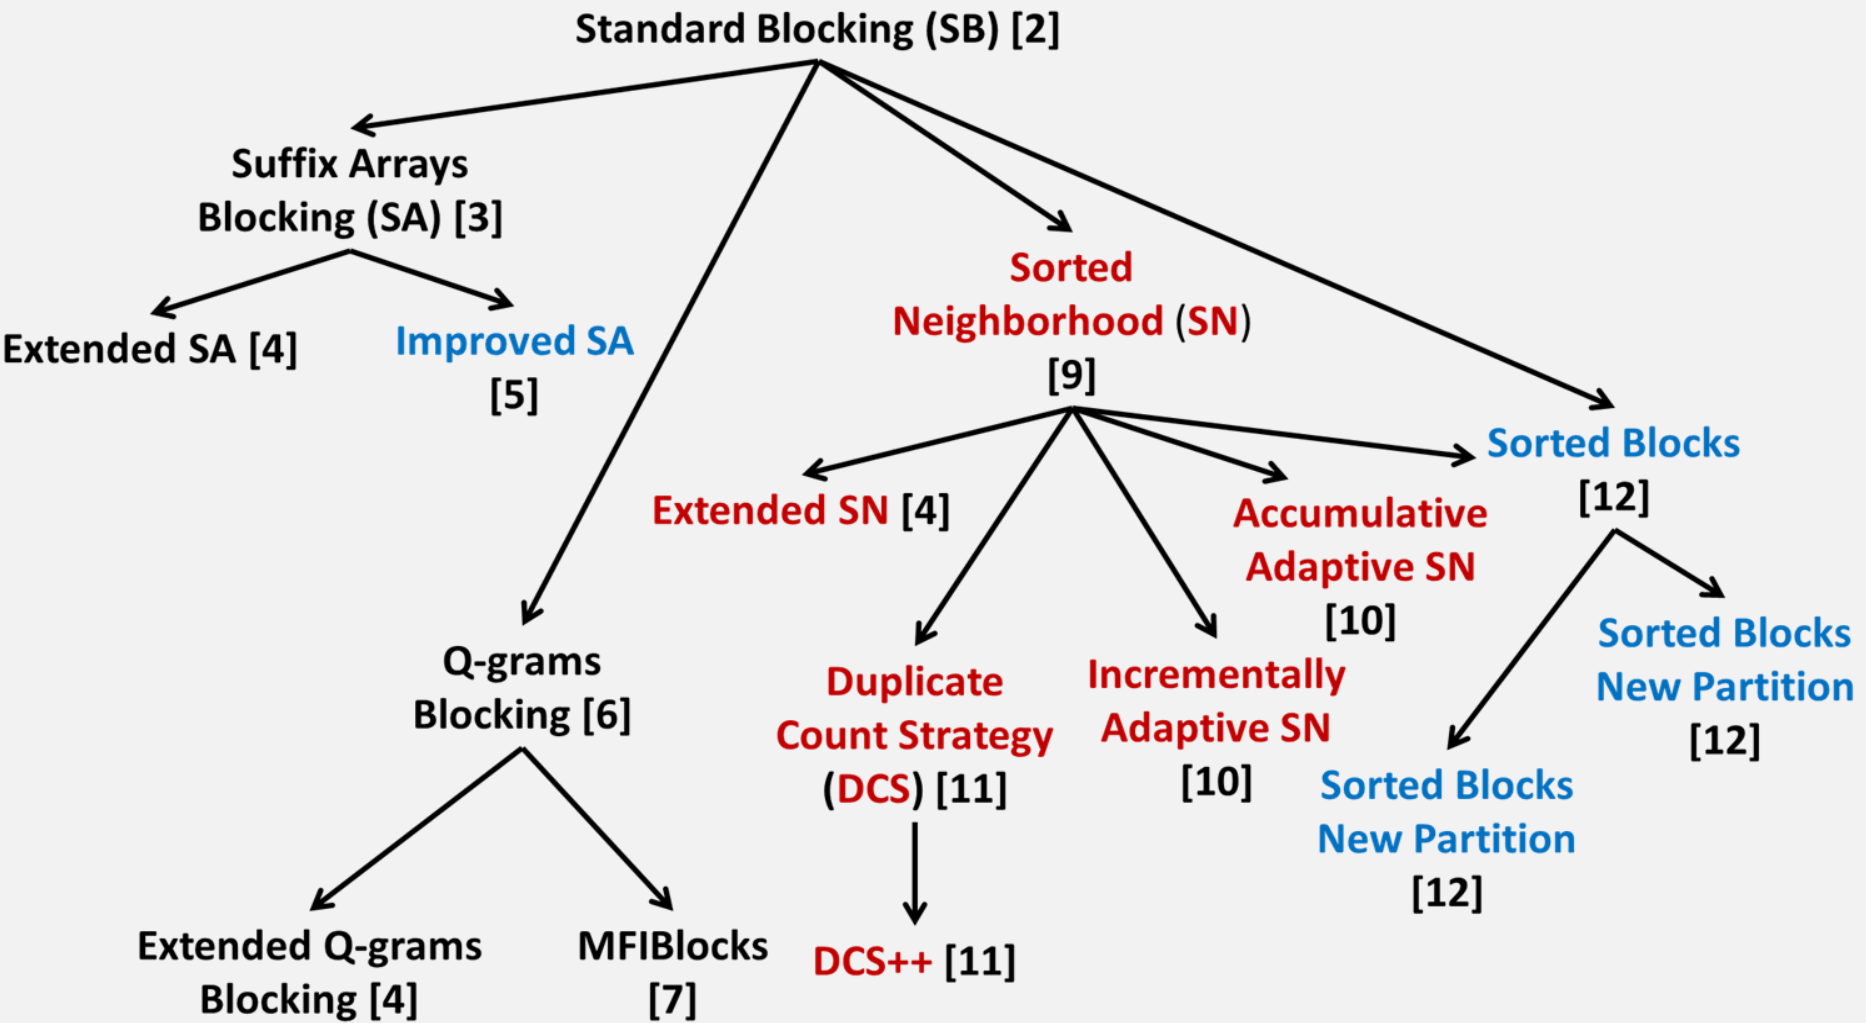
# General Principles of Blocking

1. Represent each entity by *one or more* signatures called **blocking keys**
  - Focus on **string values**
2. Place into blocks all entities having the *same or similar* blocking key
3. Two matching profiles can be **detected** as long as they co-occur in at least one block
  - **Trade-off** between recall and precision!


# Taxonomy of Blocking Methods [1]

Method	Key Type	Redundancy awareness	Matching awareness	Key selection
Standard Blocking [2]	Hash-based	Red.-free	Static	Non-learning
Suffix Arrays [3] + [4,5]	Hash-based	Red.-positive	Static	Non-learning
Q-grams Blocking [6] + [4]	Hash-based	Red.-positive	Static	Non-learning
MFIBlocks [7]	Hash-based	Red.-positive	Static	Non-learning
Sorted Neighborhood [9] + [4,10]	Sort-based	Red.-neutral	Static	Non-learning
Duplicate Count Strategy [11]	Sort-based	Red.-neutral	Dynamic	Non-learning
Sorted Blocks [12]	Hybrid	Red.-neutral	Static	Non-learning
ApproxDNF [13]	Hash-based	Red.-positive	Static	Learning-based
Blocking Scheme Learner [14]	Hash-based	Red.-positive	Static	Learning-based
CBlock [15]	Hash-based	Red.-positive	Static	Learning-based
FisherDisjunctive [16]	Hash-based	Red.-positive	Static	Learning-based

# Genealogy Tree of Non-learning Blocking Methods [1]

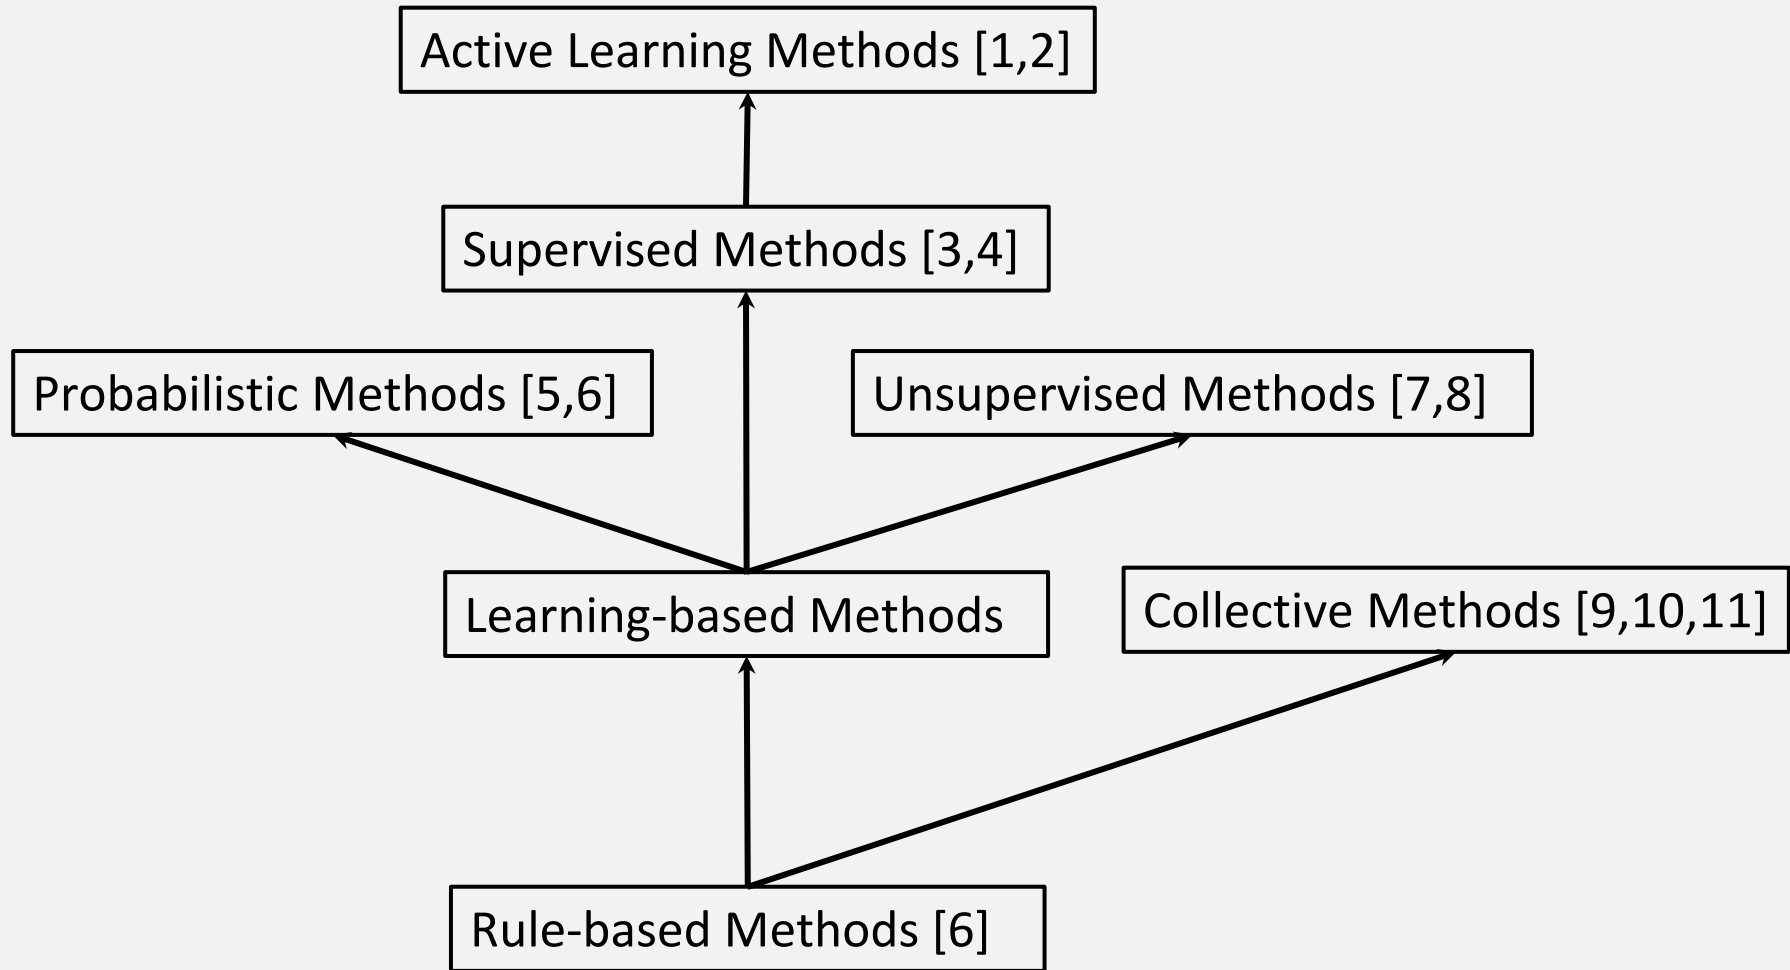


# Step 3: Matching




- Estimates the **similarity** of candidate matches.
- Input
  - A set of blocks
    - Every **distinct** comparison in any block is a candidate match
- Output
  - Similarity Graph
    - Nodes → entities
    - Edges → candidate matches
    - Edge weights → matching likelihood (based on similarity score)

# Evolution of Matching



All are heavily based on string similarity measures [6].

# Step 4: Clustering




- Partitions the matched pairs into **equivalence clusters** i.e., groups of entity profiles describing the same real-world object

- **Input**

- Similarity Graph:

- Nodes → entities
    - Edges → candidate matches
    - Edge weights → matching likelihood (based on similarity score)



- **Output**

- Equivalence Clusters

# Clustering Algorithms for Record Linkage

## Relies on **1-1 constraint**

- 1 entity from source dataset matches to 1 entity from target dataset

### 1. Unique Mapping Clustering [1][2]

- Sorts all edges in **decreasing weight**
- Starting from the top, each edge corresponds to a pair of duplicates **if**:
  - None of the adjacent entities has already been matched
  - predefined threshold < edge weight

### 2. Row-Column Clustering [3]

- efficient approximation of the **Hungarian Algorithm**

### 3. Best Assignment Clustering [4]

- efficient, heuristic solution to the **assignment problem** in unbalanced bipartite graphs

### 4. Exact Clustering [7]

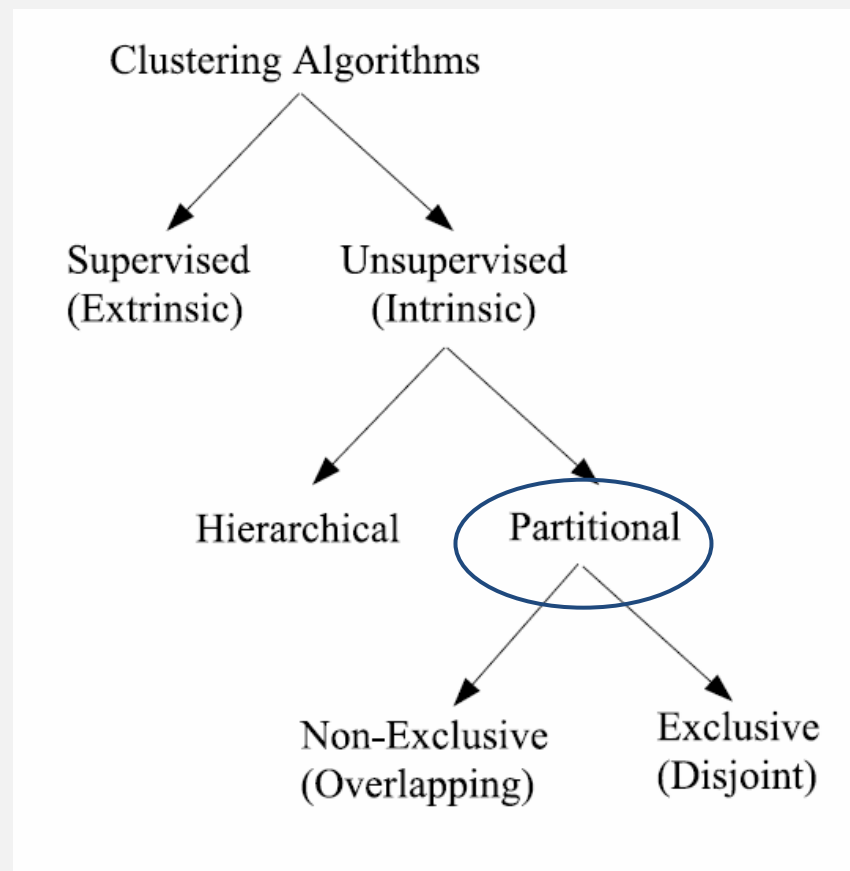
- each entity is matched with its reciprocally most similar entity

### 5. Kiraly Clustering [7]

- efficient solution to the **stable marriage problem**

# Clustering Algorithms for Deduplication

- A wealth of literature on clustering algorithms
- Requirements:
  - Partitional and disjoint Algorithms
    - Sometimes overlapping may be desirable
  - Goal: Create sets of clusters that
    - maximize the **intra-cluster** weights
    - minimize the **inter-cluster** edge weights



Classification of clustering algorithms  
[6]



# Dirty ER Clustering Algorithms Characteristics [3]

- Most important feature “**Unconstrained algorithms**”
  - Algorithms need to be able to *predict* the correct number of clusters
- Need to **scale** well
  - Time complexity  $< O(n^2)$
- Need to be **robust** with respect to characteristics of the data
  - E.g., distribution of the duplicates
- Need to be capable of finding ‘**singleton**’ clusters
  - Different from many clustering algorithms
    - E.g., algorithms proposed for image segmentation

# Summary of Experimental Results [3]

	Scalability (Current Implementations)	Ability to find the correct number of clusters	Robustness Against		
			Choice of threshold	Amount of Errors	Distribution of errors
Partitioning	High	Low	Low	Low	High
CENTER	High	High	Low	Low	High
MERGE CENTER	High	High	Low	Low	High
Star	Medium	High	Low	Low	High
SR	Low	Medium	High	High	Low
BSR	Low	Low	High	High	Low
CR	Low	High	Medium	High	High
OCR	Low	High	Medium	High	Low
Correlation Clustering	Low	High	Low	Low	High
Markov Clustering	High	High	Medium	Medium	High
Cut Clustering	Low	Low	Low	Low	High
Articulation Point	High	Medium	Low	Low	High

# Schema Matching References

1. J. Madhavan, P. A. Bernstein, and E. Rahm. Generic schema matching with cupid. In VLDB, pages 49–58, 2001.
2. S. Melnik, H. Garcia-Molina, and E. Rahm. Similarity flooding: A versatile graph matching algorithm and its application to schema matching. In ICDE, pages 117–128, 2002.
3. H.-H. Do and E. Rahm. COMA: a system for flexible combination of schema matching approaches. In VLDB, pages 610–621, 2002.
4. M. Zhang, M. Hadjieleftheriou, B. C. Ooi, C. M. Procopiuc, D. Srivastava. Automatic discovery of attributes in relational databases. In SIGMOD, pages 109–120, 2011.
5. H. W. Kuhn. The hungarian method for the assignment problem. Naval research logistics quarterly, 2(1-2):83–97, 1955.
6. L. Ramshaw and R. E. Tarjan. On minimum-cost assignments in unbalanced bipartite graphs. HP Labs, Palo Alto, CA, USA, Tech. Rep. HPL-2012-40R1, 2012.

# Blocking References – Part I

1. George Papadakis, Dimitrios Skoutas, Emmanouil Thanos, Themis Palpanas: A Survey of Blocking and Filtering Techniques for Entity Resolution. CoRR abs/1905.06167 (2019)
2. I. P. Fellegi and A. B. Sunter. A theory for record linkage. *Journal of the American Statistical Association*, 64(328):1183–1210, 1969.
3. A. N. Aizawa and K. Oyama. A fast linkage detection scheme for multi-source information integration. In *WIRI*, pages 30–39, 2005.
4. P. Christen. A survey of indexing techniques for scalable record linkage and deduplication. *IEEE TKDE*, 24(9):1537–1555, 2012.
5. T. de Vries, H. Ke, S. Chawla, and P. Christen. Robust record linkage blocking using suffix arrays. In *CIKM*, pages 305–314, 2009
6. R. Baxter, P. Christen, and T. Churches. A comparison of fast blocking methods for record linkage. In *KDD Workshops*, 2003.
7. B. Kenig and A. Gal. Mfiblocks: An effective blocking algorithm for entity resolution. *Inf. Syst.*, 38(6):908–926, 2013.
8. M. A. Hernández and S. J. Stolfo. The merge/purge problem for large databases. In *SIGMOD*, pages 127–138, 1995.
9. S. Yan, D. Lee, M. Kan, and C. L. Giles. Adaptive sorted neighborhood methods for efficient record linkage. In *JCDL*, pages 185–194, 2007.

# Blocking References – Part II

11. U. Draisbach, F. Naumann, S. Szott, and O. Wonneberg. Adaptive windows for duplicate detection. In ICDE, pages 1073–1083, 2012.
12. U. Draisbach and F. Naumann. A generalization of blocking and windowing algorithms for duplicate detection. In ICDKE, pages 18–24, 2011
13. M. Bilenko, B. Kamath, and R. J. Mooney. Adaptive blocking: Learning to scale up record linkage. In ICDM, pages 87–96, 2006
14. M. Michelson and C. A. Knoblock. Learning blocking schemes for record linkage. In AAAI, pages 440–445, 2006
15. A. D. Sarma, A. Jain, A. Machanavajjhala, and P. Bohannon. An automatic blocking mechanism for large-scale de-duplication tasks. In CIKM, pages 1055–1064, 2012.
16. M. Kejriwal and D. P. Miranker. An unsupervised algorithm for learning blocking schemes. In ICDM, pages 340–349, 2013.


# Matching References

1. K. Qian, L. Popa, P. Sen. Active Learning for Large-Scale Entity Resolution. CIKM 2017: 1379-1388
2. J. Fisher, P. Christen, Q. Wang. Active Learning Based Entity Resolution Using Markov Logic. PAKDD (2) 2016: 338-349
3. Reyes-Galaviz, O.F., Pedrycz, W., He, Z., Pizzi, N.J. A supervised gradient-based learning algorithm for optimized entity resolution. Data Knowl. Eng. 112, 106–129 (2017)
4. P. Christen. Automatic record linkage using seeded nearest neighbour and support vector machine classification. KDD 2008: 151-159.
5. A. Rasch, R. Schulze, W. Gorus, J. Hiller, S. Bartholomäus, S. Gorlatch. High-performance probabilistic record linkage via multi-dimensional homomorphisms. SAC 2019: 526-533.
6. A. K. Elmagarmid, P. G. Ipeirotis, V. S. Verykios. Duplicate Record Detection: A Survey. IEEE Trans. Knowl. Data Eng. 19(1): 1-16 (2007)
7. A. Jurek, J. Hong, Y. Chi, W. Liu. A novel ensemble learning approach to unsupervised record linkage. Inf. Syst. 71: 40-54 (2017)
8. A. Jurek, Deepak P. It Pays to Be Certain: Unsupervised Record Linkage via Ambiguity Minimization. PAKDD (3) 2018: 177-190.X
9. X. Dong, A. Y. Halevy, J. Madhavan. Reference Reconciliation in Complex Information Spaces. SIGMOD Conference 2005: 85-96.O
10. O. Benjelloun, H. Garcia-Molina, D. Menestrina, Q. Su, S. E. Whang, J. Widom. Swoosh: a generic approach to entity resolution. VLDB J. 18(1): 255-276 (2009).
11. I. Bhattacharya, L. Getoor. Collective entity resolution in relational data. TKDD 1(1): 5 (2007).

# Clustering References

1. S. Lacoste-Julien, K. Palla, A. Davies, G. Kasneci, T. Graepel, Z. Ghahramani. SIGMa: simple greedy matching for aligning large knowledge bases. KDD 2013: 572-580
2. F. M. Suchanek, S. Abiteboul, P. Senellart. PARIS: Probabilistic Alignment of Relations, Instances, and Schema. PVLDB 5(3): 157-168 (2011)
3. O. Hassanzadeh, F. Chiang, R. J. Miller, H. C. Lee. Framework for Evaluating Clustering Algorithms in Duplicate Detection. PVLDB 2(1): 1282-1293 (2009)
4. H. W. Kuhn. The hungarian method for the assignment problem. Naval research logistics quarterly, 2(1-2):83–97, 1955.
5. L. Ramshaw and R. E. Tarjan. On minimum-cost assignments in unbalanced bipartite graphs. HP Labs, Palo Alto, CA, USA, Tech. Rep. HPL-2012-40R1, 2012.
6. A. Jain, R. Dubes, “Algorithms for Clustering Data”, Prentice Hall, 1988.
7. George Papadakis, Vasilis Efthymiou, Emmanouil Thanos, Oktie Hassanzadeh, Peter Christen. An analysis of one-to-one matching algorithms for entity resolution. VLDB J. 32(6): 1369-1400 (2023)

# Generation 2: Tackling **Volume** and Veracity



- Same workflow as Generation 1
- Scope:
  - (tens of) millions of structured entity profiles
- Goals:
  - High accuracy despite noise
  - High time efficiency despite the size of data
- Assumptions:
  - Known schema → custom, schema-based solutions



# Solution: Parallelization

---

Two types:

- Multi-core parallelization
  - Single system → shared memory
  - Distribute processing among available CPUs
- Massive parallelization
  - Cluster of independent systems
  - **Map-Reduce** paradigm [1]
    - Data partitioned across the nodes of a cluster
    - Fault-tolerant, optimized execution
    - **Map Phase**: transforms a data partition into (key, value) pairs
    - **Reduce Phase**: processes pairs with the same key


# Parallelization Methods per Step

- Blocking
  - Dedoop [2]
  - MapReduce-based Sorted Neighborhood [3]
- Matching
  - Multi-core approaches [7][8]
  - MapReduce-based: Emphasis on **load balancing**
    - BlockSplit & PairRange [4][5]
    - Dis-Dedup [6]
    - Message-passing framework [9]
- Clustering
  - Fast Multi-source Entity Resolution (FAMER) framework [10][11]

# Generation 2 References

1. J. Dean and S. Ghemawat. Mapreduce: simplified data processing on large clusters. *Commun. ACM*, 51(1):107–113, 2008.
2. L. Kolb, A. Thor, and E. Rahm. Dedoop: Efficient deduplication with hadoop. *PVLDB*, 5(12):1878–1881, 2012.
3. L. Kolb, A. Thor, and E. Rahm. Multi-pass sorted neighborhood blocking with mapreduce. *Computer Science - R&D*, 27(1):45–63, 2012.
4. L. Kolb, A. Thor, and E. Rahm. Load balancing for mapreduce-based entity resolution. In *ICDE*, pages 618–629, 2012.
5. W. Yan, Y. Xue, and B. Malin. Scalable load balancing for mapreduce-based record linkage. In *IPCCC*, pages 1–10, 2013.
6. X. Chu, I. F. Ilyas, and P. Koutris. Distributed data deduplication. *PVLDB*, 9(11):864–875, 2016.
7. O. Benjelloun, H. Garcia-Molina, H. Gong, H. Kawai, T. E. Larson, D. Menestrina, and S. Thavisomboon. D-swoosh: A family of algorithms for generic, distributed entity resolution. In *ICDCS*, page 37, 2007.
8. Hung-sik Kim and Dongwon Lee. Parallel linkage. In *CIKM*, pages 283–292, 2007.
9. V. Rastogi, N. N. Dalvi, and M. N. Garofalakis. Large-scale collective entity matching. *PVLDB*, 4(4):208–218, 2011.
10. A. Saeedi, E. Peukert, and E. Rahm. Comparative evaluation of distributed clustering schemes for multi-source entity resolution. In *ADBIS*, pages 278–293, 2017.
11. A. Saeedi, M. Nentwig, E. Peukert, and E. Rahm. Scalable matching and clustering of entities with FAMER. *CSIMQ*, 16:61–83, 2018.

# G3: Tackling **Variety**, Volume and Veracity



- Scope:

- User-generated Web Data

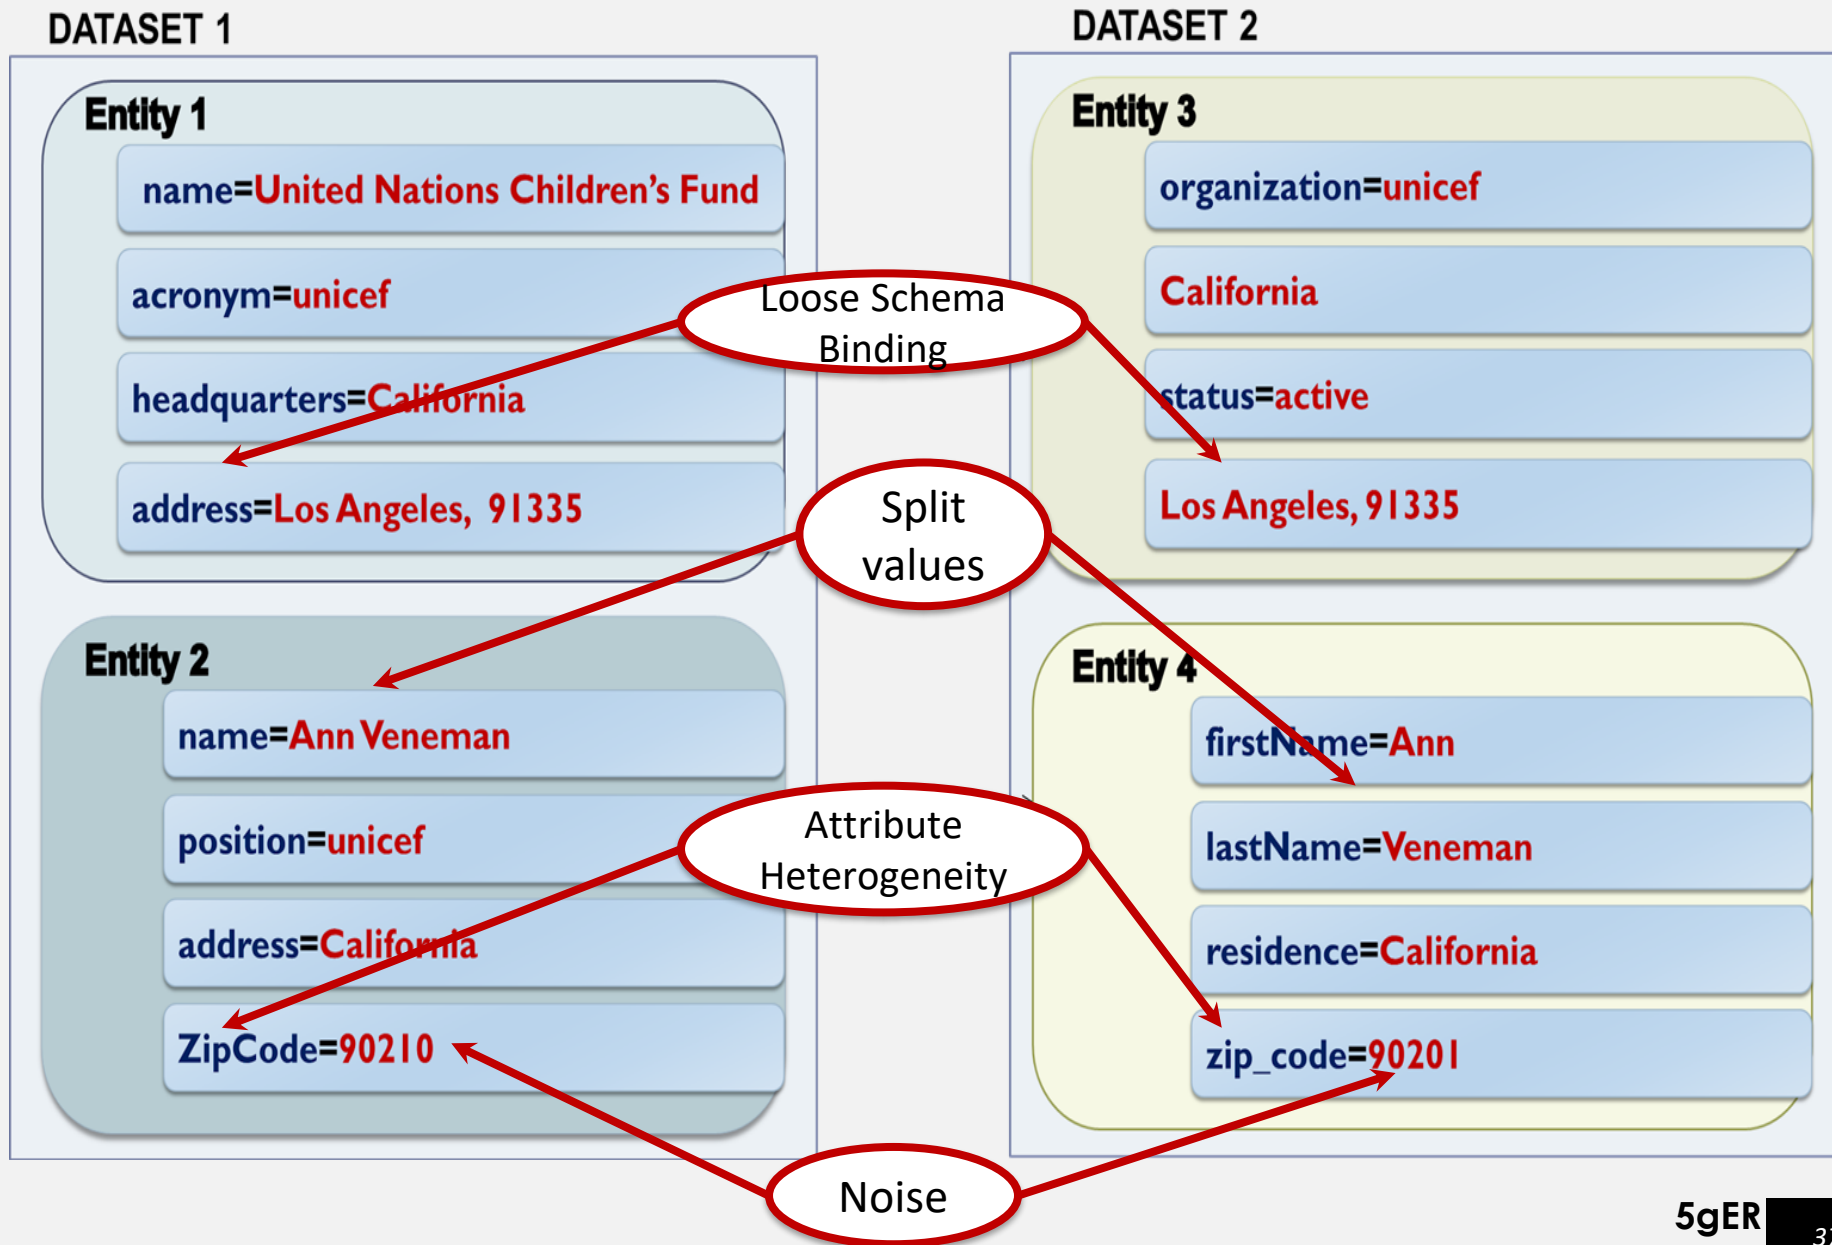
Voluminous, **(semi-)structured datasets**.

- BTC09: **1.15 billion** triples, **182 million** entities


**Users are free to add attribute values** and/or attribute names  
unprecedented levels of schema heterogeneity.

- Google Base: **100,000** schemata for **10,000** entity types
- BTC09: **136,000** attribute names

# Example of Web Data




# Schema Clustering



- Schema Matching → **not applicable**
- Instead, partition attributes according to their **syntactic** similarity, regardless of their **semantic** relation
- Goal:
  - Facilitate next steps
- Scope:
  - Both Clean-Clean and Dirty ER
- Attribute Clustering [1][2][3]
  - Create a graph, where every node represents an attribute
  - For each attribute name/node  $n_i$ 
    - Find the most similar node  $n_j$
    - If  **$\text{sim}(n_i, n_j) > 0$** , add an edge  $\langle n_i, n_j \rangle$
  - Extract connected components
  - Put all singleton nodes in a **“glue” cluster**

# Block Building



- Unlike Blocking in G1/G2, it considers **all** attribute **values** and completely ignores all attribute names  
→ **schema-agnostic functionality**
- Core approach: **Token Blocking** [1]
  1. Given an entity profile, extract all tokens that are contained in its attribute values.
  2. Create one block for every distinct token with frequency  $> 2$  → each block contains all entities with the corresponding token.

Pros:

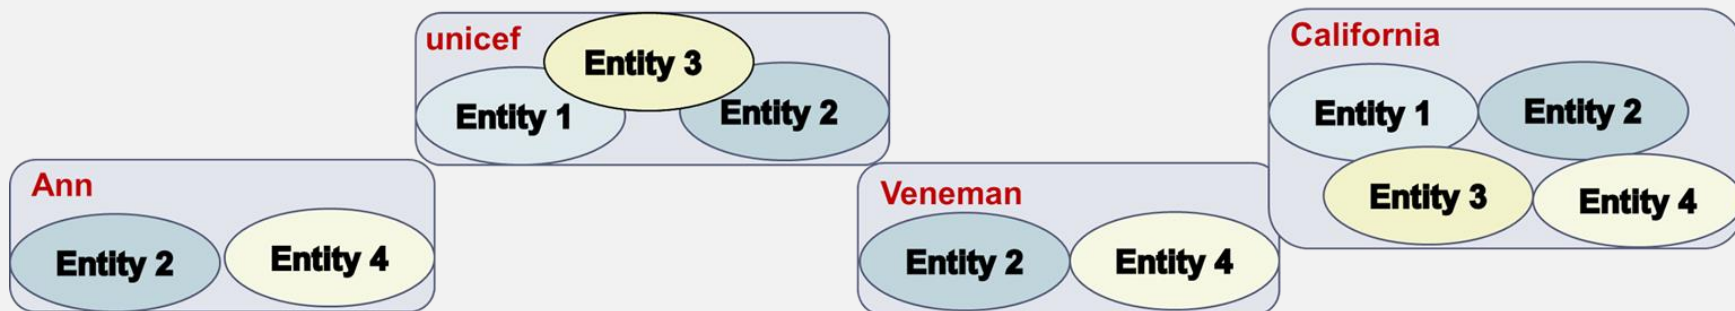
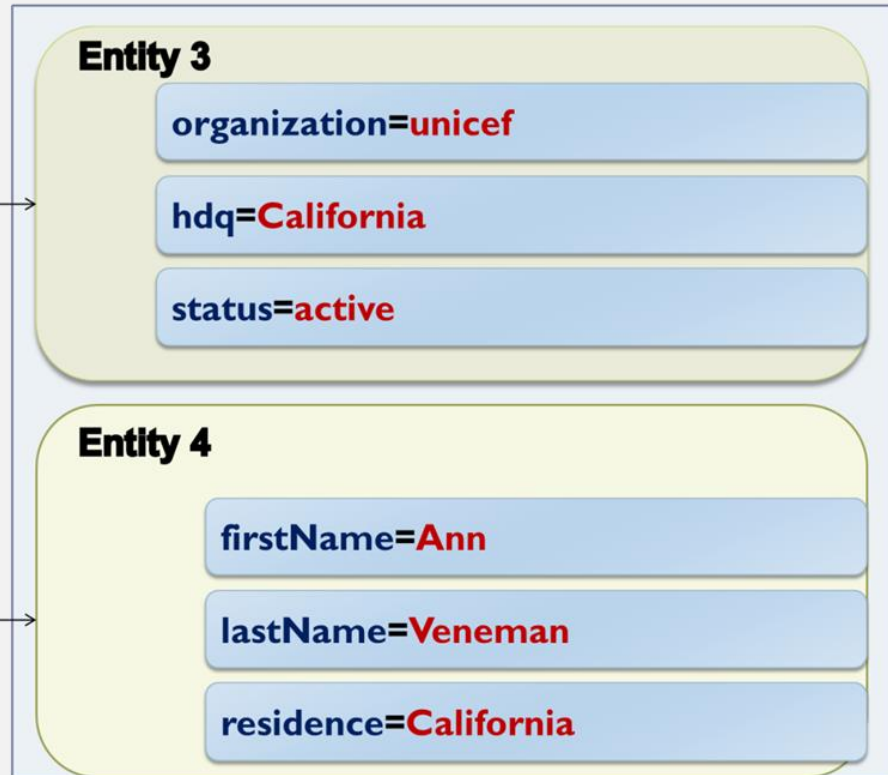
  - Parameter-free
  - Efficient
  - Unsupervised

# Example of Token Blocking

DATASET 1

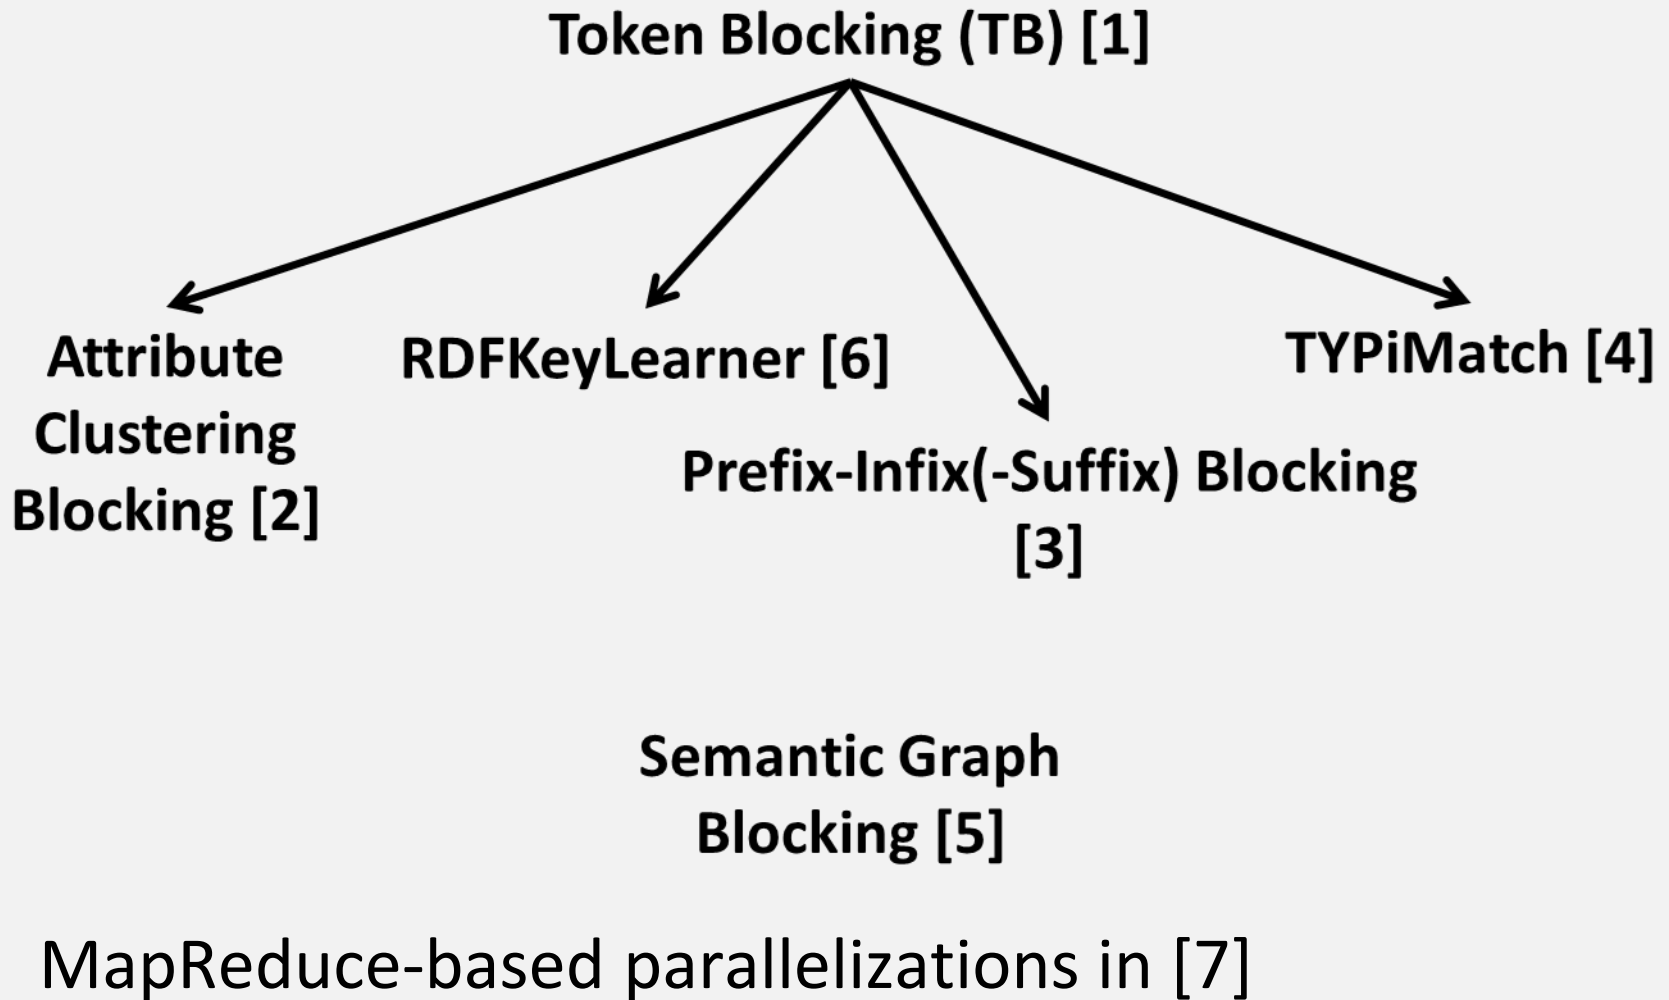


DATASET 2






# Genealogy of Block Building Techniques [8]



# Block Processing



- High **Recall** due to redundancy
- Low **Precision** due to:
  1. the blocks are overlapping → **redundant comparisons**
  2. high number of comparisons between irrelevant entities → **superfluous comparisons**

## Solution:

restructure the original blocks so as to increase **precision** at no significant cost in **recall**

# Block Processing Techniques

---

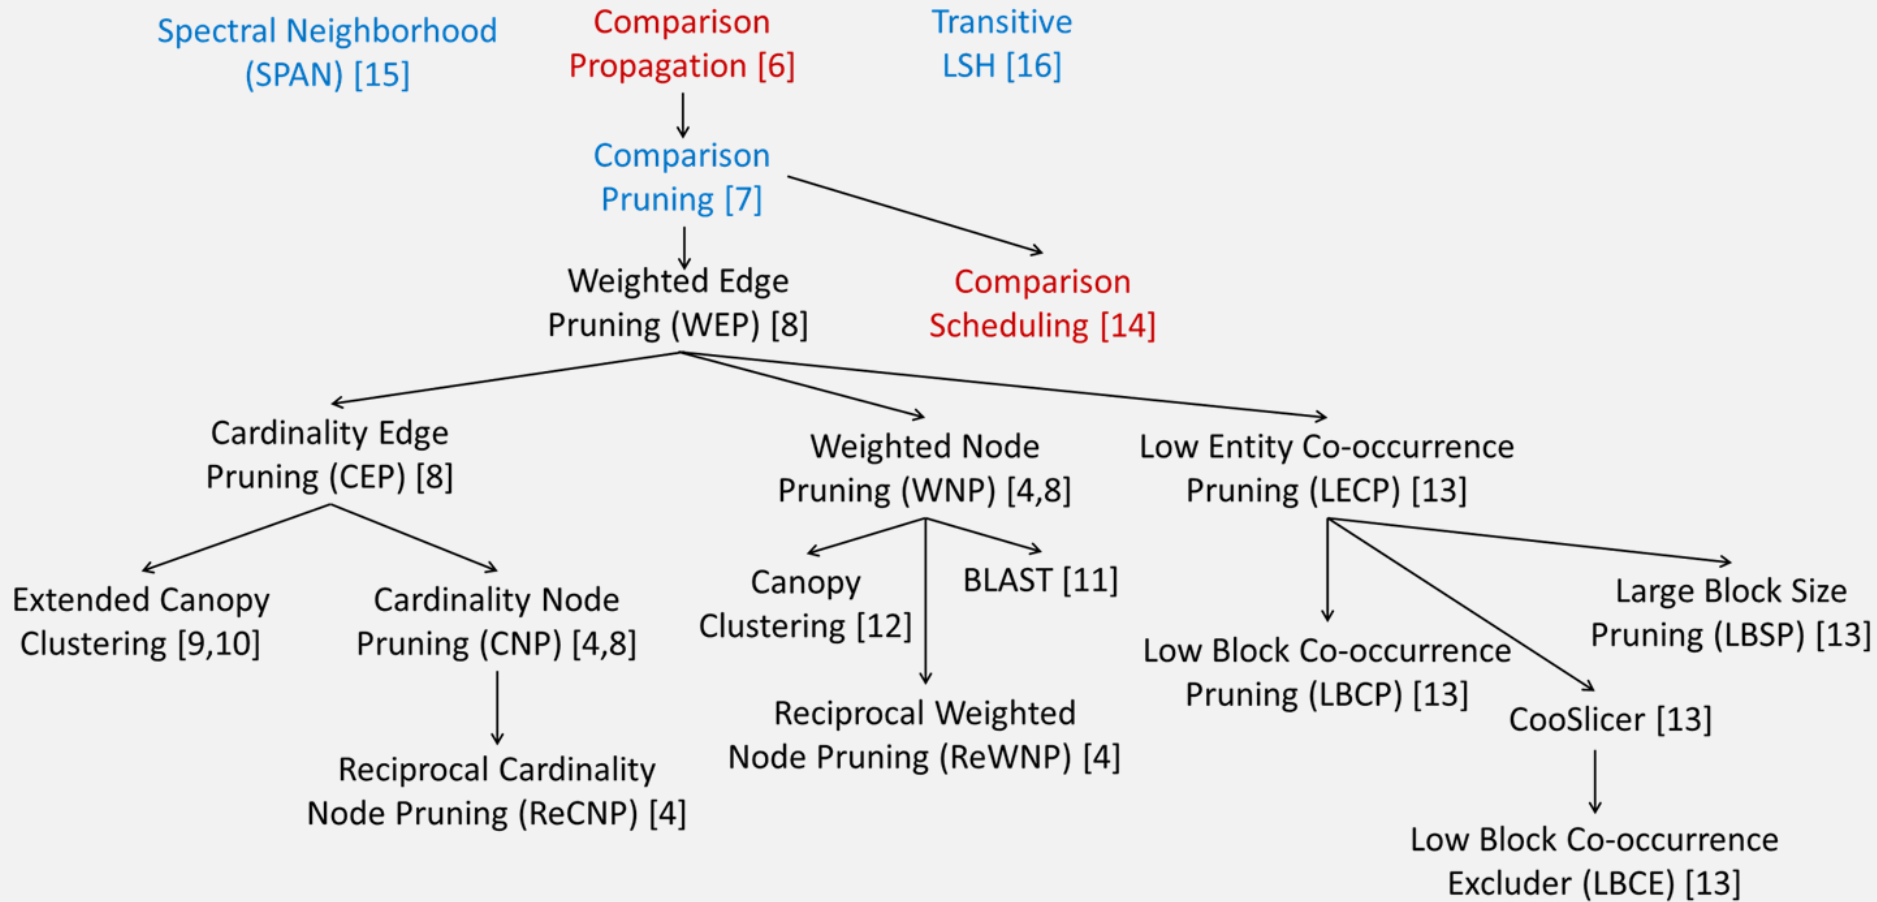
## Generic approach

- Assign a **matching likelihood score** to each item
- Discard items with low costs


## Block-centric methods

- Block Purging [1,2,3]
- Block Filtering [4]
- Block Clustering [5]


# Comparison Cleaning Methods [17]



# Entity Matching



- Collective approaches to **tackle Variety**
- Most methods crafted for **Clean-Clean ER**
- General outline of SiGMa [1], PARIS [2], LINDA [3], RiMOM-IM [4,5]
  - Bootstrap with a few **reliable seed** matches.
  - Using value and neighbor similarity, propagate initial matches to neighbors.
  - Order candidate matches in **descending** overall similarity
  - Iteratively mark the **top pair** as a match if it satisfies a constraint
  - Recompute the similarity of the neighbors
  - Update candidate matches order
- MinoanER [6] performs a specific number of steps, rather than iterating until convergence



- Methods of G1 & G2 are **still applicable**
  - Only difference: similarity scores extracted in a schema-agnostic fashion, not from specific attributes
- SplitMerge [1]
  - inherently capable of handling heterogeneous semantic types

[1] M. Nentwig, A. Groß, and E. Rahm. Holistic entity clustering for linked data. In ICDM Workshops, pages 194–201, 2016.

# Schema Clustering References

1. G. Papadakis, E. Ioannou, T. Palpanas, C. Niederée, W. Nejdl. A Blocking Framework for Entity Resolution in Highly Heterogeneous Information Spaces. *IEEE Trans. Knowl. Data Eng.* 25(12): 2665-2682 (2013)
2. G. Simonini, S. Bergamaschi, H. V. Jagadish. BLAST: a Loosely Schema-aware Meta-blocking Approach for Entity Resolution. *PVLDB* 9(12): 1173-1184 (2016)
3. G. Papadakis, L. Tsekouras, E. Thanos, G. Giannakopoulos, T. Palpanas, M. Koubarakis. The return of JedAI: End-to-End Entity Resolution for Structured and Semi-Structured Data. *PVLDB* 11(12): 1950-1953 (2018)

# Block Building References

1. G. Papadakis, E. Ioannou, C. Niederée, P. Fankhauser. Efficient entity resolution for large heterogeneous information spaces. WSDM 2011: 535-544
2. G. Papadakis, E. Ioannou, T. Palpanas, C. Niederée, W. Nejdl. A Blocking Framework for Entity Resolution in Highly Heterogeneous Information Spaces. IEEE Trans. Knowl. Data Eng. 25(12): 2665-2682 (2013)
3. G. Papadakis, E. Ioannou, C. Niederée, T. Palpanas, W. Nejdl. Beyond 100 million entities: large-scale blocking-based resolution for heterogeneous data. WSDM 2012: 53-62
4. Y. Ma, T. Tran. TYPiMatch: type-specific unsupervised learning of keys and key values for heterogeneous web data integration. WSDM 2013: 325-334
5. J. Nin, V. Muntés-Mulero, N. Martínez-Bazan, and J. Larriba-Pey. On the use of semantic blocking techniques for data cleansing and integration. In IDEAS, pages 190–198, 2007.
6. D. Song and J. Heflin. Automatically generating data linkages using a domain-independent candidate selection approach. In ISWC, pages 649–664, 2011.
7. V. Christophides, V. Efthymiou, K. Stefanidis. Entity Resolution in the Web of Data. Synthesis Lectures on the Semantic Web: Theory and Technology, Morgan & Claypool Publishers 2015.
8. George Papadakis, Dimitrios Skoutas, Emmanouil Thanos, Themis Palpanas: A Survey of Blocking and Filtering Techniques for Entity Resolution. CoRR abs/1905.06167 (2019)



# Block Processing References – Part I

1. G. Papadakis, E. Ioannou, C. Niederée, P. Fankhauser. Efficient entity resolution for large heterogeneous information spaces. WSDM 2011: 535-544
2. G. Papadakis, E. Ioannou, T. Palpanas, C. Niederée, W. Nejdl. A Blocking Framework for Entity Resolution in Highly Heterogeneous Information Spaces. IEEE Trans. Knowl. Data Eng. 25(12): 2665-2682 (2013)
3. G. Papadakis, E. Ioannou, C. Niederée, T. Palpanas, W. Nejdl. Beyond 100 million entities: large-scale blocking-based resolution for heterogeneous data. WSDM 2012: 53-62
4. G. Papadakis, G. Papastefanatos, T. Palpanas, M. Koubarakis. Scaling Entity Resolution to Large, Heterogeneous Data with Enhanced Meta-blocking. EDBT 2016: 221-232
5. J. Fisher, P. Christen, Q. Wang, E. Rahm. A Clustering-Based Framework to Control Block Sizes for Entity Resolution. KDD 2015: 279-288
6. G. Papadakis, E. Ioannou, C. Niederée, T. Palpanas, W. Nejdl. Eliminating the redundancy in blocking-based entity resolution methods. JCDL 2011: 85-94.
7. G. Papadakis, E. Ioannou, C. Niederée, T. Palpanas, W. Nejdl. To compare or not to compare: making entity resolution more efficient. SWIM 2011: 3.
8. G. Papadakis, G. Koutrika, T. Palpanas, W. Nejdl. Meta-Blocking: Taking Entity Resolution to the Next Level. IEEE Trans. Knowl. Data Eng. 26(8): 1946-1960 (2014).
9. P. Christen. A Survey of Indexing Techniques for Scalable Record Linkage and Deduplication. IEEE Trans. Knowl. Data Eng. 24(9): 1537-1555 (2012).
10. G. Papadakis, G. Alexiou, G. Papastefanatos, G. Koutrika. Schema-agnostic vs Schema-based Configurations for Blocking Methods on Homogeneous Data. PVLDB 9(4): 312-323 (2015).

# Block Processing References – Part II

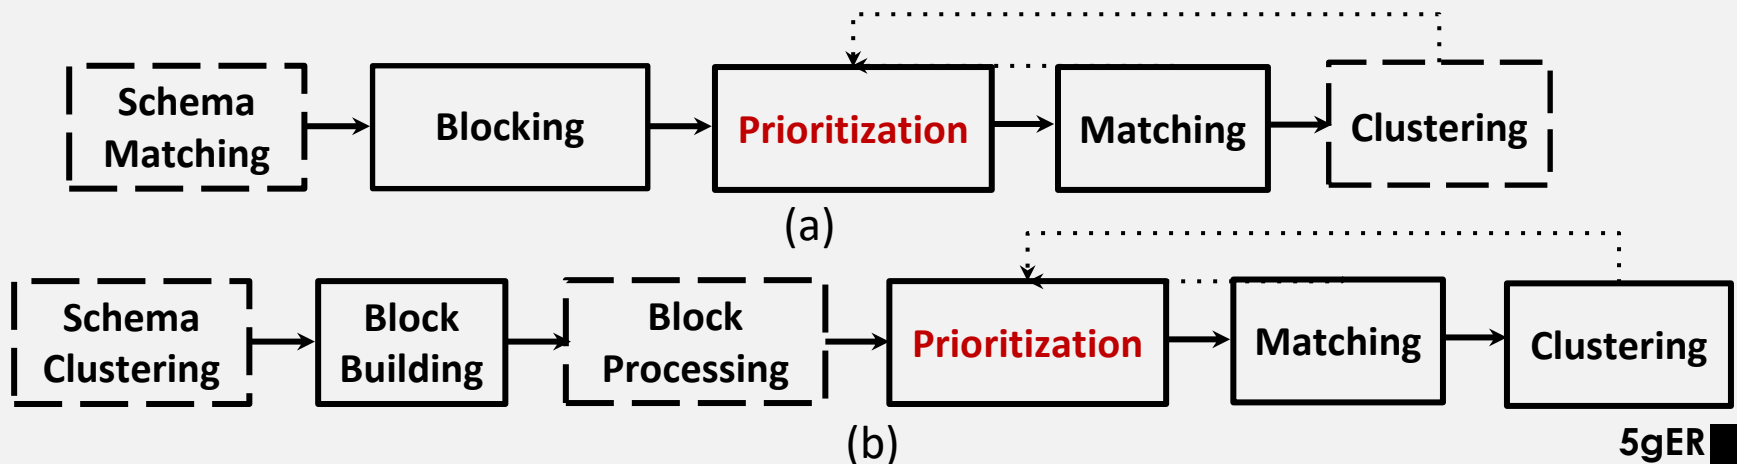
11. G. Simonini, S. Bergamaschi, H. V. Jagadish. BLAST: a Loosely Schema-aware Meta-blocking Approach for Entity Resolution. PVLDB 9(12): 1173-1184 (2016)
12. A. McCallum, K. Nigam, L. H. Ungar. Efficient clustering of high-dimensional data sets with application to reference matching. KDD 2000: 169-178.
13. D. C. Nascimento, C. E. S. Pires, and D. G. Mestre. Exploiting block co-occurrence to control block sizes for entity resolution. Knowledge and Information Systems, pages 1–42, 2019.
14. G. Papadakis, E. Ioannou, T. Palpanas, C. Niederée, and W. Nejdl. A blocking framework for entity resolution in highly heterogeneous information spaces. IEEE TKDE, 25(12):2665–2682, 2013.
15. L. Shu, A. Chen, M. Xiong, and W. Meng. Efficient spectral neighborhood blocking for entity resolution. In ICDE, pages 1067–1078, 2011.
16. R. C. Steorts, S. L. Ventura, M. Sadinle, and S. E. Fienberg. A comparison of blocking methods for record linkage. In Privacy in Statistical Databases, pages 253–268, 2014.
17. George Papadakis, Dimitrios Skoutas, Emmanouil Thanos, Themis Palpanas: A Survey of Blocking and Filtering Techniques for Entity Resolution. CoRR abs/1905.06167 (2019)

# Entity Matching References

1. S. Lacoste-Julien, K. Palla, A. Davies, G. Kasneci, T. Graepel, Z. Ghahramani. SIGMa: simple greedy matching for aligning large knowledge bases. KDD 2013: 572-580
2. F. M. Suchanek, S. Abiteboul, P. Senellart. PARIS: Probabilistic Alignment of Relations, Instances, and Schema. PVLDB 5(3): 157-168 (2011)
3. C. Böhm, G. de Melo, F. Naumann, and G. Weikum. LINDA: distributed web-of-data-scale entity matching. In CIKM, pages 2104–2108, 2012.
4. J. Li, J. Tang, Y. Li, and Q. Luo. Rimom: A dynamic multistrategy ontology alignment framework. TKDE, 21(8):1218–1232, 2009.
5. C. Shao, L. Hu, J. Li, Z. Wang, T. L. Chung, and J.-B. Xia. Rimom-im: A novel iterative framework for instance matching. J. Comput. Sci. Technol., 31(1):185–197, 2016.
6. V. Efthymiou, G. Papadakis, K. Stefanidis, and V. Christophides. MinoanER: Schema-agnostic, non-iterative, massively parallel resolution of web entities. In EDBT, pages 373–384, 2019.

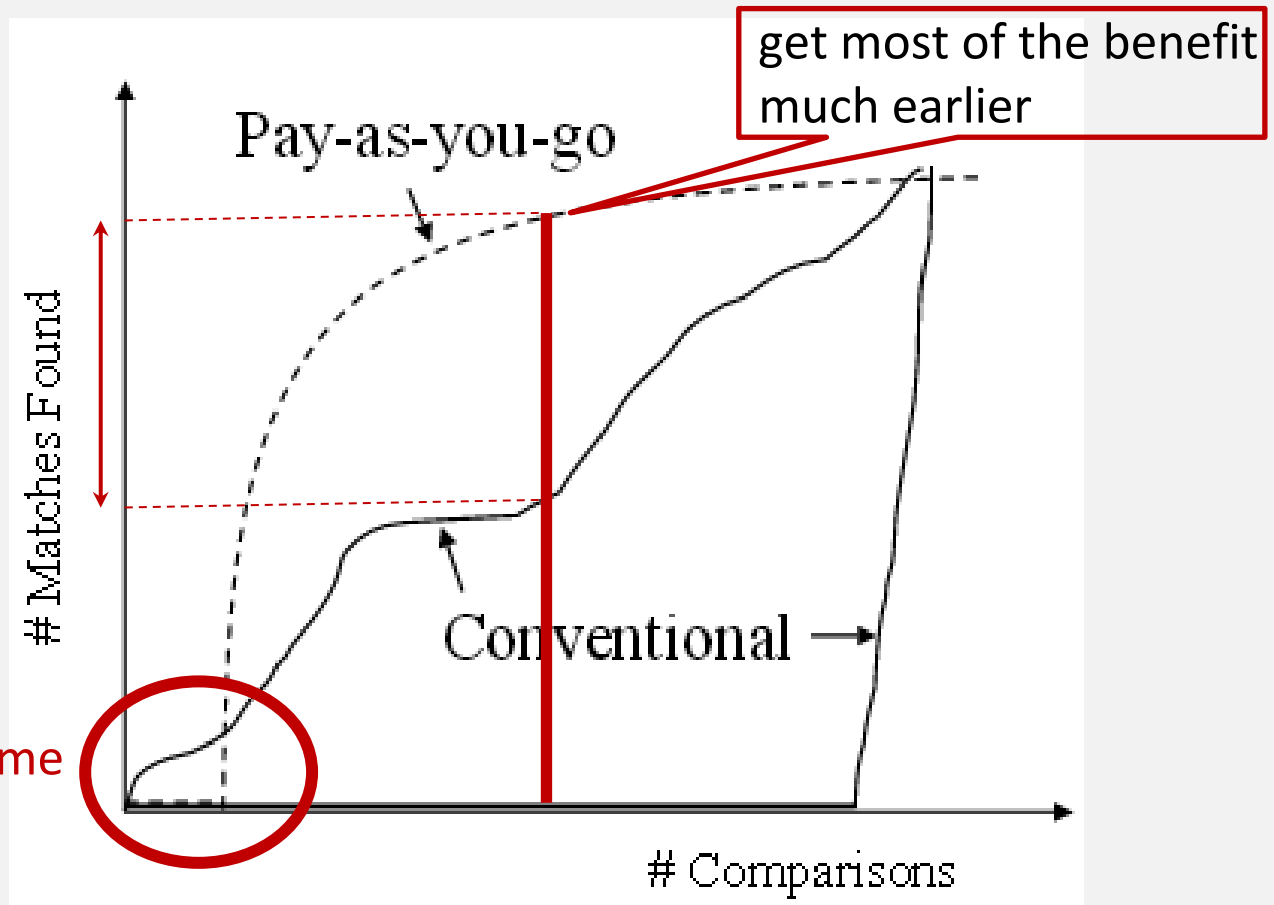
# G4: Tackling **Velocity**, Variety, Volume and Veracity

- Scope:
  - Applications with increasing data volume and time constraints
    - Loose ones (e.g., minutes, hours) → **Progressive ER**
    - Strict ones (i.e., seconds) → **Real-time (On-line) ER**
- End-to-end workflows for Progressive ER



# Progressive Entity Resolution

Unprecedented, increasing volume of data → applications requiring partial solutions to produce useful results

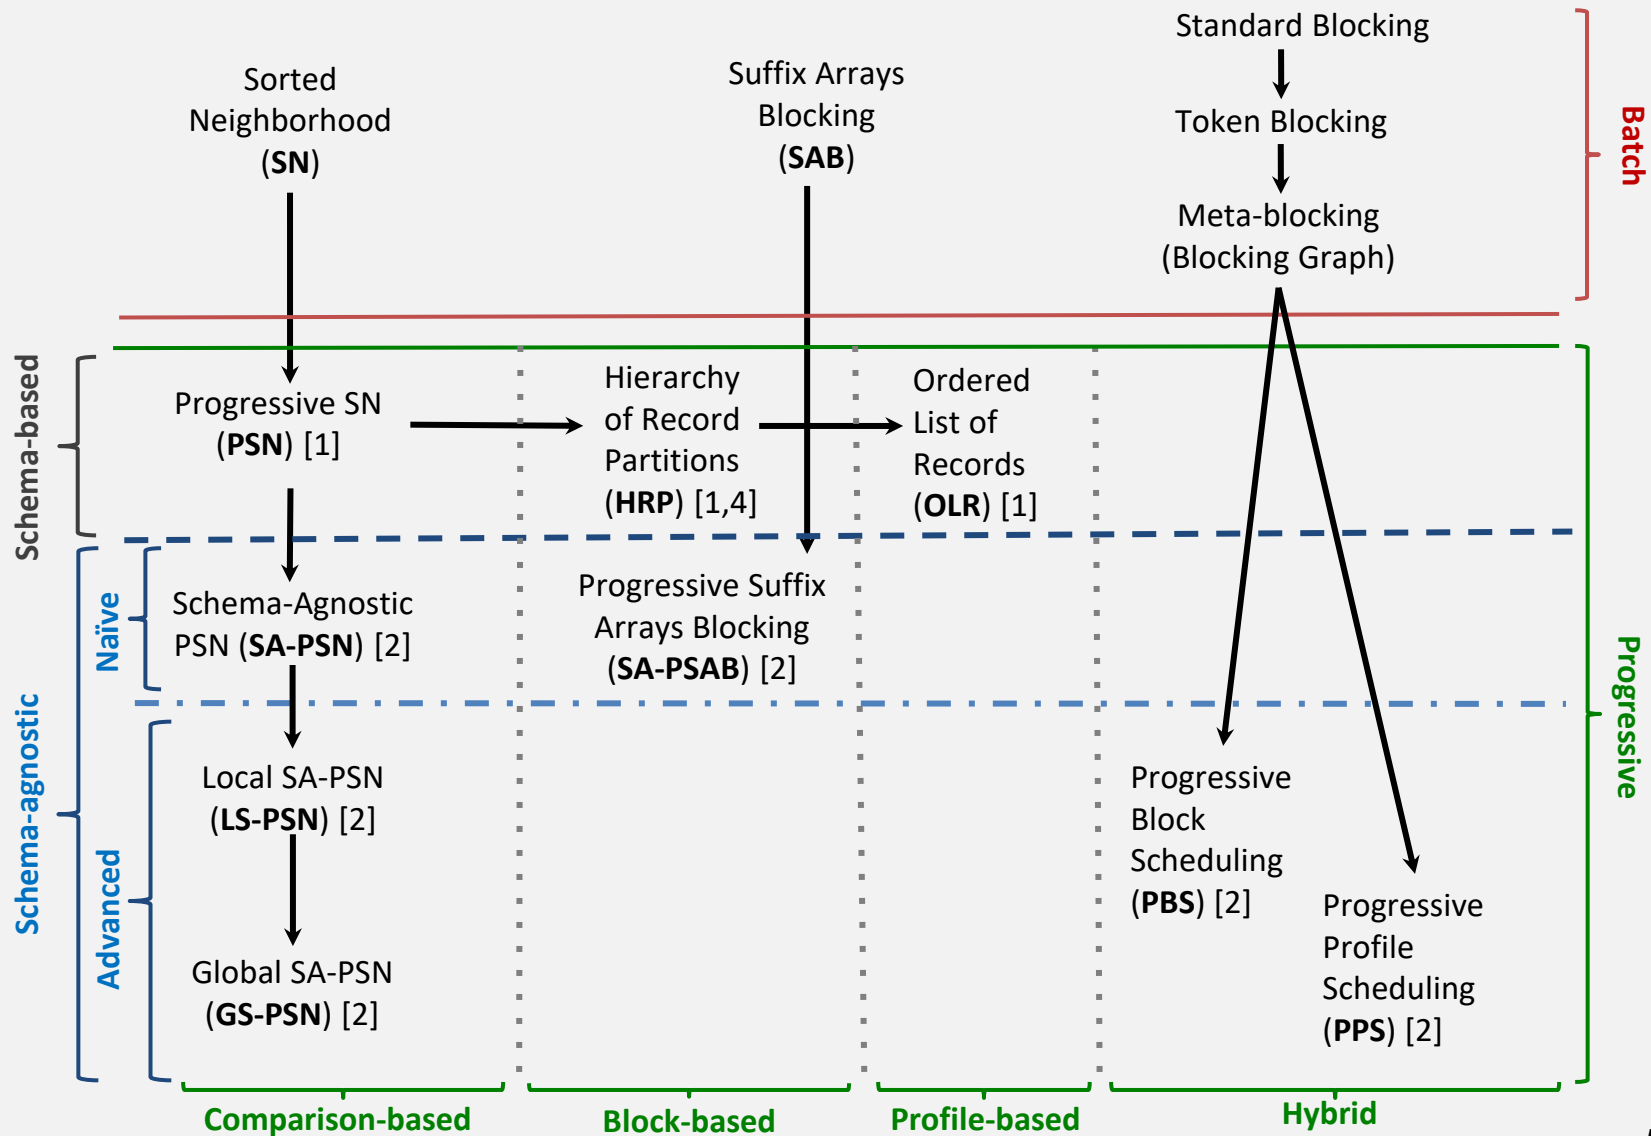


may require some pre-processing

# Outline Progressive ER


- Requires:
  - Improved Early Quality
  - Same Eventual Quality
- **Prioritization**
  - Defines **optimal processing order** for a set of entities
  - Static Methods [1,2]:
    - Guide which records to compare first, **independently** of Entity Matching results
  - Dynamic Methods [3]:
    - If  $c_{i,j}$  is a duplicate, then check  $c_{i+1,j}$  and  $c_{i,j+1}$  as well.
    - Assumption:
      - Oracle for Entity Matching

# Taxonomy of Static Prioritization Methods



# Real-time Entity Resolution

**Same** workflow as Generations 1 and 2:



**Same** scope (so far):

- Structured data

**Different** input:

- **stream** of query entity profiles

**Different** goal:

- resolve each query over a large dataset in the shortest possible time (& with the minimum memory footprint)



# Techniques per workflow step

## Incremental Blocking

- **DySimII** [1] - extends Standard Blocking
- **F-DySNI** [2,3] - extends Sorted Neighborhood
- **(S)BlockSketch** [4] - bounded matching time, constant memory footprint

## Incremental Matching

- **QDA** [5] - SQL-like selection queries over a single dataset
- **QuERy** [6] - complex join queries over multiple, overlapping, dirty DSs
- **EAQP** [7] - queries under data
- Evolving matching rules [8]

## Incremental Clustering

- Incremental Correlation Clustering [9]

# Progressive ER References

1. S. E. Whang, D. Marmaros, and H. Garcia-Molina. Pay-as-you-go entity resolution. *TKDE*, 25(5):1111–1124, 2013.
2. T. Papenbrock, A. Heise, and F. Naumann. Progressive duplicate detection. *TKDE*, 27(5):1316–1329, 2015.
3. G. Simonini, G. Papadakis, T. Palpanas, S. Bergamaschi. Schema-Agnostic Progressive Entity Resolution. *IEEE Trans. Knowl. Data Eng.* 31(6): 1208-1221 (2019)
4. Y. Altowim and S. Mehrotra. Parallel progressive approach to entity resolution using mapreduce. In *ICDE*, pages 909–920, 2017.

# Incremental ER References

1. B. Ramadan and P. Christen, H. Liang, and R. W. Gayler, and D. Hawking. Dynamic similarity-aware inverted indexing for real-time entity resolution. In PAKDD Workshops, pages 47–58, 2013.
2. B. Ramadan and P. Christen. Forest-based dynamic sorted neighborhood indexing for real-time entity resolution. In CIKM, pages 1787–1790, 2014.
3. B. Ramadan and P. Christen, H. Liang, and R. W. Gayler. Dynamic sorted neighborhood indexing for real-time entity resolution. *J. Data and Information Quality*, 6(4):15:1–15:29, 2015.
4. D. Karapiperis, A. Gkoulalas-Divanis, V. S. Verykios. Summarization Algorithms for Record Linkage. *EDBT 2018*: 73-84.
5. H. Altwaijry, D. V. Kalashnikov, and S. Mehrotra. QDA: A query-driven approach to entity resolution. *TKDE*, 29(2):402–417, 2017.
6. H. Altwaijry, S. Mehrotra, and D. V. Kalashnikov. Query: A framework for integrating entity resolution with query processing. *PVLDB*, 9(3):120–131, 2015.
7. E. Ioannou, W. Nejdl, C. Niederée, and Y. Velegrakis. On-the-fly entity-aware query processing in the presence of linkage. *PVLDB*, 3(1): 429–438, 2010.
8. S. E. Whang and H. Garcia-Molina. Entity resolution with evolving rules. *PVLDB*, 3(1):1326–1337, 2010.
9. A. Gruenheid, X. L. Dong, and D. Srivastava. Incremental record linkage. *Proc. VLDB Endow.*, 7(9):697–708, May 2014. ISSN 2150-8097.

# G5: Leveraging External Knowledge

- Applies to any of the previous generations
  - **No change** in the end-to-end workflows
- Performance improves by incorporating one of the three types of contextual information:
  1. **Human** common sense through **crowd-sourcing**
    - Open web data through:
  2. **Pre-trained** Language Models (**PLMs**)
  3. **Large** Language Models (**LLMs**)
- PLMs apply to **both** blocking and matching, unlike crowd-sourcing and LLMs, which apply **exclusively** to matching

# Crowd-sourcing

- Process/work divided among a large number of people, either paid or unpaid
- Idea: tasks are **simple** for human intelligence, but **complex** for computers
- Approach:
  - Break a problem into microtasks, called Human Intelligence Tasks (**HITS**)
  - Choose an online community
    - [Amazon Mechanical Turk](#)
    - [Figure Eight](#) (former CrowdFlower)
  - Assign to every individual, called **worker**, a series of HITs
  - Each worker is paid per executed HIT → **monetary cost**
  - Popular for solving many tasks, e.g., CrowdDB

# Crowd-sourcing for Entity Resolution

- Delegate the **entity matching decisions** to the workers i.e., transform pairwise comparisons into HITs
- Challenges:
  1. **Generating HITs:** CrowdER [8], ZenCrowd [9]
  2. **Formulating HITs:**  
Pair- & cluster-based [8], Hybrid [10], Crowdlink [14]
  3. **Balancing accuracy and monetary cost:**  
Random ordering [3], probabilistic question selection [2], Edge- and node-centric ordering [1], maximize progressive recall [4], adaptive crowd-based deduplication [12], attribute labeling and clustering [15], partial-order based framework [17], bDENSE [18], probabilistic ER with crowd errors [11, 16], and pair-wise error correction layer [13]
  4. **Restricting the labor cost:**  
Corleone [5], Falcon [6], and CloudMatcher [7]

# Crowd-sourced ER References – Part I

1. J. Wang, G. Li, T. Kraska, M. J. Franklin, and J. Feng. Leveraging transitive relations for crowdsourced joins. In SIGMOD, pages 229–240, 2013.
2. S. E. Whang, P. Lofgren, and H. Garcia-Molina. Question selection for crowd entity resolution. PVLDB, 6(6):349–360, 2013.
3. N. Vesdapunt, K. Bellare, and N. N. Dalvi. Crowdsourcing algorithms for entity resolution. PVLDB, 7(12):1071–1082, 2014.
4. D. Firmani, B. Saha, and D. Srivastava. Online entity resolution using an oracle. PVLDB, 9(5):384–395, 2016.
5. C. Gokhale, S. Das, AnHai Doan, J. F. Naughton, Narasimhan Rampalli, Jude W. Shavlik, Xiaojin Zhu. Corleone: hands-off crowdsourcing for entity matching. SIGMOD Conference 2014: 601-6122.
6. S. Das, P. Suganthan G. C., AnHai Doan, J. F. Naughton, G. Krishnan, R. Deep, E. Arcaute, V. Raghavendra, Y.Park. Falcon: Scaling Up Hands-Off Crowdsourced Entity Matching to Build Cloud Services. SIGMOD Conference 2017: 1431-14463.
7. Y. Govind, E. Paulson, P. Nagarajan, P. Suganthan G. C., AnHai Doan, Y. Park, G. Fung, D. Conathan, M. Carter, M. Sun. CloudMatcher: A Hands-Off Cloud/Crowd Service for Entity Matching. PVLDB 11(12): 2042-2045 (2018).
8. Jiannan Wang, Tim Kraska, Michael J. Franklin, Jianhua Feng. CrowdER: Crowdsourcing Entity Resolution. PVLDB 5(11): 1483-1494 (2012).
9. Gianluca Demartini, Djellel Eddine Difallah, Philippe Cudré-Mauroux. ZenCrowd: leveraging probabilistic reasoning and crowdsourcing techniques for large-scale entity linking. WWW 2012: 469-478.

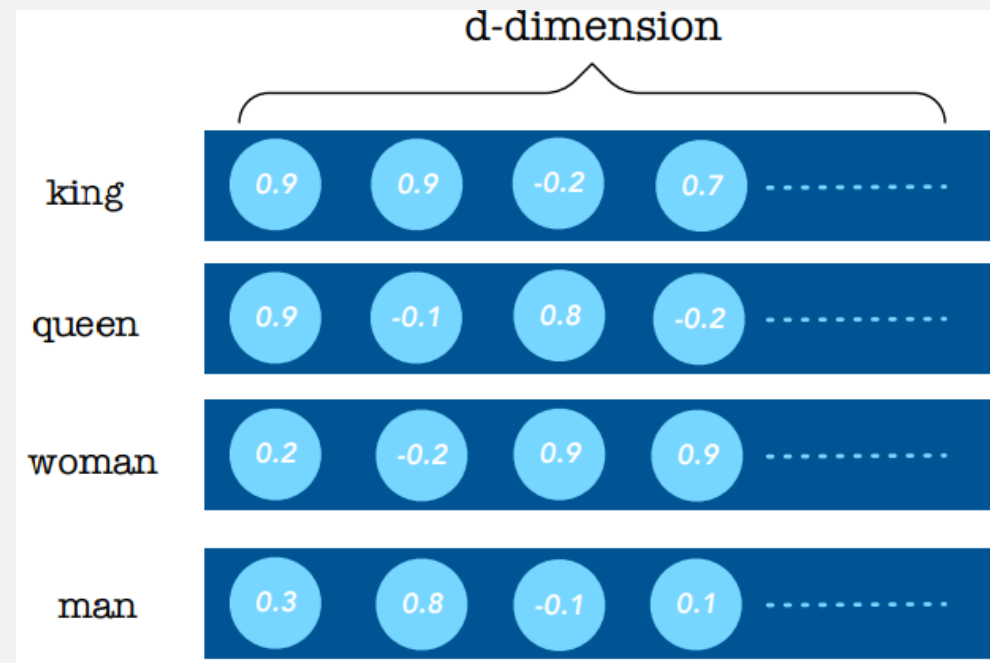
# Crowd-sourced ER References – Part II

10. Vasilis Verroios, Hector Garcia-Molina, and Yannis Papakonstantinou. Waldo: An adaptive human interface for crowd entity resolution. In SIGMOD, pages 1133–1148, 2017.
11. V. K. Yalavarthi, X. Ke, and A. Khan. Select Your Questions Wisely: For Entity Resolution with Crowd Errors. In CIKM, pages 317-326, 2017.
12. S. Wang, X. Xiao, and C.-H. Lee. Crowd-Based Deduplication: An Adaptive Approach. In SIGMOD, pages 1263-1277, 2015.
13. S. Galhotra, D. Firmani, B. Saha, and D. Srivastava. Robust entity resolution using random graphs. In SIGMOD, pages 3–18, 2018.
14. C. J. Zhang, R. Meng, L. Chen, and F. Zhu. Crowmlink: An error-tolerant model for linking complex records. In ExploreDB, pages 15–20, 2015.
15. A. R. Khan and H. Garcia-Molina. Attribute-based crowd entity resolution. In CIKM, pages 549–558, 2016.
16. X. Ke, M. Teo, A. Khan, V. K. Yalavarthi. A Demonstration of PERC: Probabilistic Entity Resolution With Crowd Errors. PVLDB 11(12): 1922-1925 (2018)
17. C. Chai, G. Li, J. Li, D. Deng, and J. Feng. Cost-effective crowdsourced entity resolution: A partial-order approach. In SIGMOD, pages 969–984, 2016.
18. V. Verroios and H. Garcia-Molina. Entity resolution with crowd errors. In ICDE, pages 219–230, 2015.




# Embeddings

- Based on the **distributional hypothesis**  
i.e., words appearing in the **same context share meaning**
- Each word is represented as a distribution of weights (positive or negative) across specific dimensions
- Goal: capture **semantic** string similarities
- Popular embeddings pre-trained over huge corpora:
  - Word2Vec [5]
  - Glove [6]
  - fastText [7]



# Deep Learning

- Specific class of Machine Learning / Data Mining
- Teaches computers to do what comes naturally to humans: learn by example
- Goal: learn a complicated function from the data
- Ideal for **complex** tasks involving **multi-dimensional** data like the embedding vectors of **PLMs**
- Has transformed many fields, e.g., computer vision, speech recognition, natural language processing, etc.
  - Similar performance, or even better, to human expert performance
- Details in [1]



# Initial Approaches of Deep Learning

- SEMPROP [2] for **schema matching**
  - Semantic + syntactic matcher
- AutoBlock [3] for **blocking**
  - Combines similarity-preserving representation learning with nearest neighbor search
- DeepMatcher [8], Multi-Perspective Matching [9], and DeepER [4] for **matching**
  - Attribute embedding, summarization, and comparison
  - Deep Learning solutions
- Following approaches
  - Improve weaknesses

Architecture module		Options
Attribute embedding		<i>Granularity:</i> (1) Word-based (2) Character-based
		<i>Training:</i> (3) Pre-trained (4) Learned
Attribute similarity representation	(1) Attribute summarization	(1) Heuristic-based (2) RNN-based (3) Attention-based (4) Hybrid
	(2) Attribute comparison	(1) Fixed distance (cosine, Euclidean) (2) Learnable distance (concatenation, element-wise absolute difference, element-wise multiplication)
Classifier		NN (multi-layer perceptron)

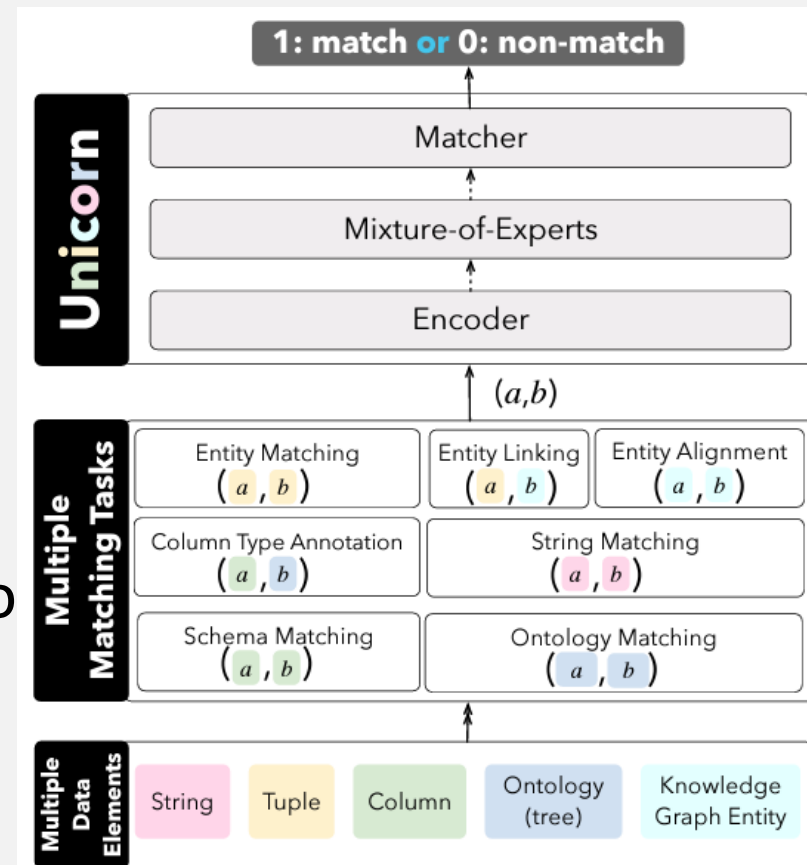
# HierGAT [10]

---

- Weaknesses of existing initial approaches
  - Assume all words / attributes are equally important
  - Don't consider that words from different domains may have different meanings
- Create and process resolution using a **Graph**
- Encodes entities, attributes, and words
- Captures related relationships
- Assigns different weights given category

# Unicorn [11]

- Weaknesses of existing initial approaches:
  - Task-specific solutions that disable the opportunities for generalization or sharing learnt knowledge
- Proposed a unified model for “data matching” task in data integration
  - Encoder: converts pair  $(a,b)$  into a learned representation
  - Mixture-of-Experts: enhances the learned representation into a better representation
  - Matcher: binary classifier



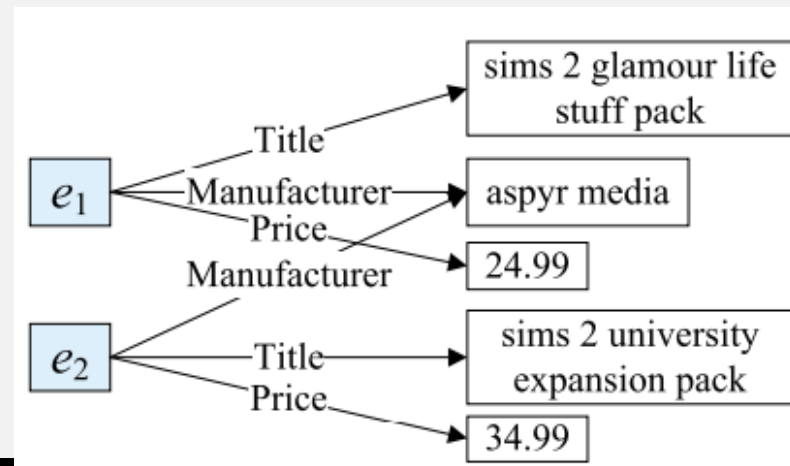
# Sudowoodo [12]

---

- Weaknesses of existing initial approaches:
  - Require creating large-scale, high quality labeled datasets
  - Require separate modeling, annotation, and experimentation for each (sub) task of the process
- Contrastive learning: self-supervision approach that learns data representations where similar data items are close while different ones are far apart
  - Done by pre-training a representation model
- This fine-tuned model is used to generate the embeddings.
- The learned representations either directly used or facilitate fine-tuning to support different tasks.

# CollaborEM [13]

- Weaknesses of existing initial approaches
  - Require a large number of labeled pairs
  - Insufficient feature discovery
- Generate labeled tuple pairs construct a graph that
  - Is the smallest, i.e., with fewer nodes and edges than graphs of other approaches
  - Preserves the semantic relationships between each tuple and its corresponding attribute values and between different tuples via shared value-level nodes



# Deep Learning References

1. Ian J. Goodfellow, Yoshua Bengio, Aaron C. Courville. Deep Learning. Adaptive computation and machine learning, MIT Press 2016, ISBN 978-0-262-03561-3, pp. 1-775
2. R. C. Fernandez, E. Mansour, A. A. Qahtan, A. K. Elmagarmid, I. F. Ilyas, S. Madden, M. Ouzzani, M. Stonebraker, N. Tang. Seeping Semantics: Linking Datasets Using Word Embeddings for Data Discovery. ICDE 2018: 989-1000
3. W. Zhang, H. Wei, B. Sisman, X. L. Dong, C. Faloutsos, D. Page. AutoBlock: A Hands-off Blocking Framework for Entity Matching. WSDM 2020: 744-752
4. M. Ebraheem, S. Thirumuruganathan, S. R. Joty, M. Ouzzani, N. Tang. Distributed Representations of Tuples for Entity Resolution. PVLDB 11(11): 1454-1467 (2018)
5. J. Pennington, R. Socher, and C. D. Manning. Glove: Global vectors for word representation. In EMNLP, pages 1532–1543
6. T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean. Distributed representations of words and phrases and their compositionality. In NIPS, pages 3111–3119, 2013
7. P. Bojanowski, E. Grave, A. Joulin, and T. Mikolov. Enriching word vectors with subword information. TACL, 5:135–146, 2017



# Deep Learning References

8. S. Mudgal, H. Li, T. Rekatsinas, A.H. Doan, Y. Park, G. Krishnan, R. Deep, E. Arcaute, V. Raghavendra. Deep Learning for Entity Matching: A Design Space Exploration. SIGMOD Conference 2018: 19-34
9. C. Fu, X. Han, L. Sun, B. Chen, W. Zhang, S. Wu, H. Kong. End-to-End Multi-Perspective Matching for Entity Resolution. IJCAI 2019: 4961-4967
10. Dezhong Yao, Yuhong Gu, Gao Cong, Hai Jin, and Xinqiao Lv. 2022. Entity Resolution with Hierarchical Graph Attention Networks. In SIGMOD Conference. ACM, 429-442.
11. Jianhong Tu, Ju Fan, Nan Tang, Peng Wang, Guoliang Li, Xiaoyong Du, Xiaofeng Jia, and Song Gao. 2023. Unicorn: A Unified Multi-tasking Model for Supporting Matching Tasks in Data Integration. Proc. ACM Manag. Data 1, 1 (2023), 84:1-84:26.
12. Runhui Wang, Yuliang Li, and Jin Wang. 2023. Sudowoodo: Contrastive Self-supervised Learning for Multi-purpose Data Integration and Preparation. In ICDE. IEEE, 1502-1515.
13. Congcong Ge, Pengfei Wang, Lu Chen, Xiaoze Liu, Baihua Zheng, Yunjun Gao: CollaborEM: A Self-Supervised Entity Matching Framework Using Multi-Features Collaboration. IEEE Trans. Knowl. Data Eng. 35(12): 12139-12152 (2023)

# Large Language Models (LLMs)

---

- Core idea: ask a chatbot **whether a given pair of entity descriptions are matching or not**
  - each question is called “**prompt**”
- Challenge:
  - Unlike PLMs, the embedding representation is **transparent**
  - They constitute interactive approaches that are **sensitive** to the form of the **prompt**
- Solutions:
  - Prompt engineering!

# Basic prompt engineering [1]

Three parameters that can be configured **independently**:

1. Problem definition:

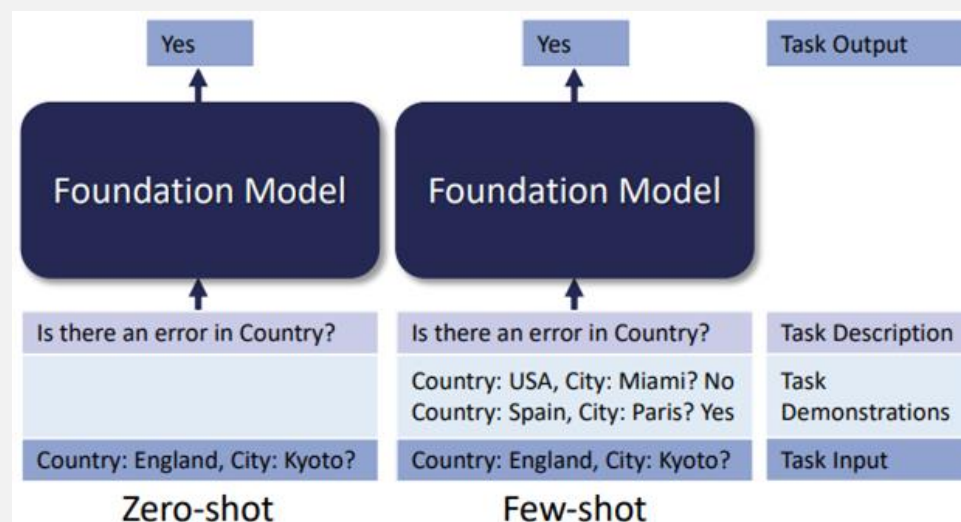
- "Are Product A and Product B the same?" **OR**
- "Are Product A and Product B equivalent?"

2. In-context learning:

- zero shot **OR**
- few shot
  - Random Selection **OR**
  - Manual Selection by experts

3. Entity serialization:

- with all attributes **OR**
- with a subset of attributes



## Conclusions for **GPT3-175B**:

- Few shot outperforms zero shot to a significant extent
- Attribute selection is better than using all attributes
- Problem definition can have a large impact
- Comparable performance with DL-based matching algorithms

# Fine-grained prompt engineering [2]

Three parameters that can be configured **independently**:

## 1. Problem definition

- General (they refer to entities)
- Domain specific (they refer to entity types, e.g., products)

## 2. Language

- Simple (e.g., do two entities match?)
- Complex (e.g., do two entities refer to the same real-world product?)

## 3. Output

- Free (no output specifications)
- Forced (e.g., reply "Yes" or "No")

## 4. Entity Serialization

- Single attribute
- Multiple attributes

## 5. In-context learning

- Zero shot
- Few shot
  - examples selection
    - at random, by expert or by context similarity
  - number of examples (e.g., 6, 10 or 20)

## 6. Instructions with matching rules

## 7. Fine-tuning

<b>Task Desc.</b>	Do the following two product descriptions match?
<b>Demonstrations</b>	Product 1: 'Title: DYMO D1 19 mm x 7 m' Product 2: 'Title: Dymo D1 (19mm x 7m – BoW)'
<b>Answer</b>	Yes.
<b>Task Desc.</b>	Do the following two product descriptions match?
<b>Demonstrations</b>	Product 1: 'Title: DYMO D1 Tape 24mm' Product 2: 'Title: Dymo D1 19mm x 7m'
<b>Answer</b>	No.
<b>Task Desc.</b>	Do the following two product descriptions match?
<b>Task Input</b>	Product 1: 'Title: DYMO D1 – Roll (1.9cm x 7m)' Product 2: 'Title: DYMO D1 Tape 12mm x 7m'

# Fine-grained prompt engineering – Part II

Conclusions using 6 LLMs:

- 3 hosted:
  1. gpt3.5-turbo-0301
  2. gpt3.5-turbo-0613
  3. gpt4-0613
- 3 open-source:
  1. SOLAR 70B
  2. Beluga2
  3. Mixtral-8x7B

Main takeaways:

1. No prompt consistently outperforms all others
2. Open-source LLMs have similar effectiveness with hosted ones
3. LLMs comparable with DL-based matchers even in zero-shot settings
4. Few shot and instruction-based prompts outperform zero shot
5. Fine-tuning significantly improves effectiveness

# Prompt strategies [3]

## Three different approaches:

### 1. Match strategy:

Pair-wise questions (as in previous works)

### 2. Comparison strategy:

Given two entities, find the most similar to a specific entity.

### 3. Selection strategy:

Given k candidates for a specific entity, identify the matching one or none of them.

**Do these two records refer to the same real-world entity?**

(1) *Cruzer Force USB Flash Drive 32GB Type-A 2.0 Chrome*  
(2) *Sandisk USB Flash Drive 32GB Cruzer Glide 2.0/3.0*

**LLM Response** No

(a) Matching

**Which of these two records is more consistent with the given record:**

*Cruzer Force USB Flash Drive 32GB Type-A 2.0 Chrome*  
(A) *Sandisk USB Flash Drive 32GB Cruzer Glide 2.0/3.0*  
(B) *Pendrive Sandisk Cruzer Force - SDCZ71-032G-B35*

**LLM Response** Record B

(b) Comparing

**Select a record from the following list that refers to the same real-world entity as the given record:**

*Cruzer Force USB Flash Drive 32GB Type-A 2.0 Chrome*  
(1) *Sandisk USB Flash Drive 32GB Cruzer Glide 2.0/3.0*  
(2) *Pendrive Sandisk Cruzer Force - SDCZ71-032G-B35*  
(3) *Sandisk Extreme Pro 3.1 Solid State Flash Drive 128GB*  
(4) *Kingston DataTraveler G4 32 GB USB-stick*

...

**LLM Response** Record 2

(c) Selecting

# Batch Prompting [4]

- Goal: **reduce the cost of hosted LLMs**, which charge in proportion to the number of input tokens, through batching, i.e., multiple pairwise questions with the same demonstrations.
- **BatchER** options:
  - Question Batching based on PLMs or structure-aware similarities like Jaccard similarity or edit distance
    - Random
    - Similarity-based (using clustering algorithms like DBScan and K-Means )
    - Diversity-based (using one pair from each similarity-based cluster)
  - Demonstration selection
    - Fixed
    - Top-k batch, i.e., the k most relevant demonstrations per batch
    - Top-k question, i.e., the most relevant demonstration per pair in the batch
    - Covering-based, i.e., for each pair in the batch, there is a demonstration with distance lower than a threshold
- Conclusions:
  - Batch prompting outperforms standard prompting both to effectiveness and cost
  - Best performance corresponds to Diversity-based Question Batching with Covering-based Demonstration Selection

# LLMs References

---

1. Avanika Narayan, Ines Chami, Laurel J. Orr, Christopher Ré. Can Foundation Models Wrangle Your Data? Proc. VLDB Endow. 16(4): 738-746 (2022)
2. Ralph Peeters, Christian Bizer. Entity Matching using Large Language Models. CoRR abs/2310.11244 (2023)
3. T Wang, H Lin, X Chen, X Han, H Wang, Z Zeng, L Sun. Match, Compare, or Select? An Investigation of Large Language Models for Entity Matching. arXiv preprint arXiv:2405.16884, 2024.
4. Meihao Fan, Xiaoyue Han, Ju Fan, Chengliang Chai, Nan Tang, Guoliang Li, Xiaoyong Du. Cost-Effective In-Context Learning for Entity Resolution: A Design Space Exploration. CoRR abs/2312.03987 (2023)



- Introduction
- Generations: 1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup>, 5<sup>th</sup>

## Part C: Hands-on Session

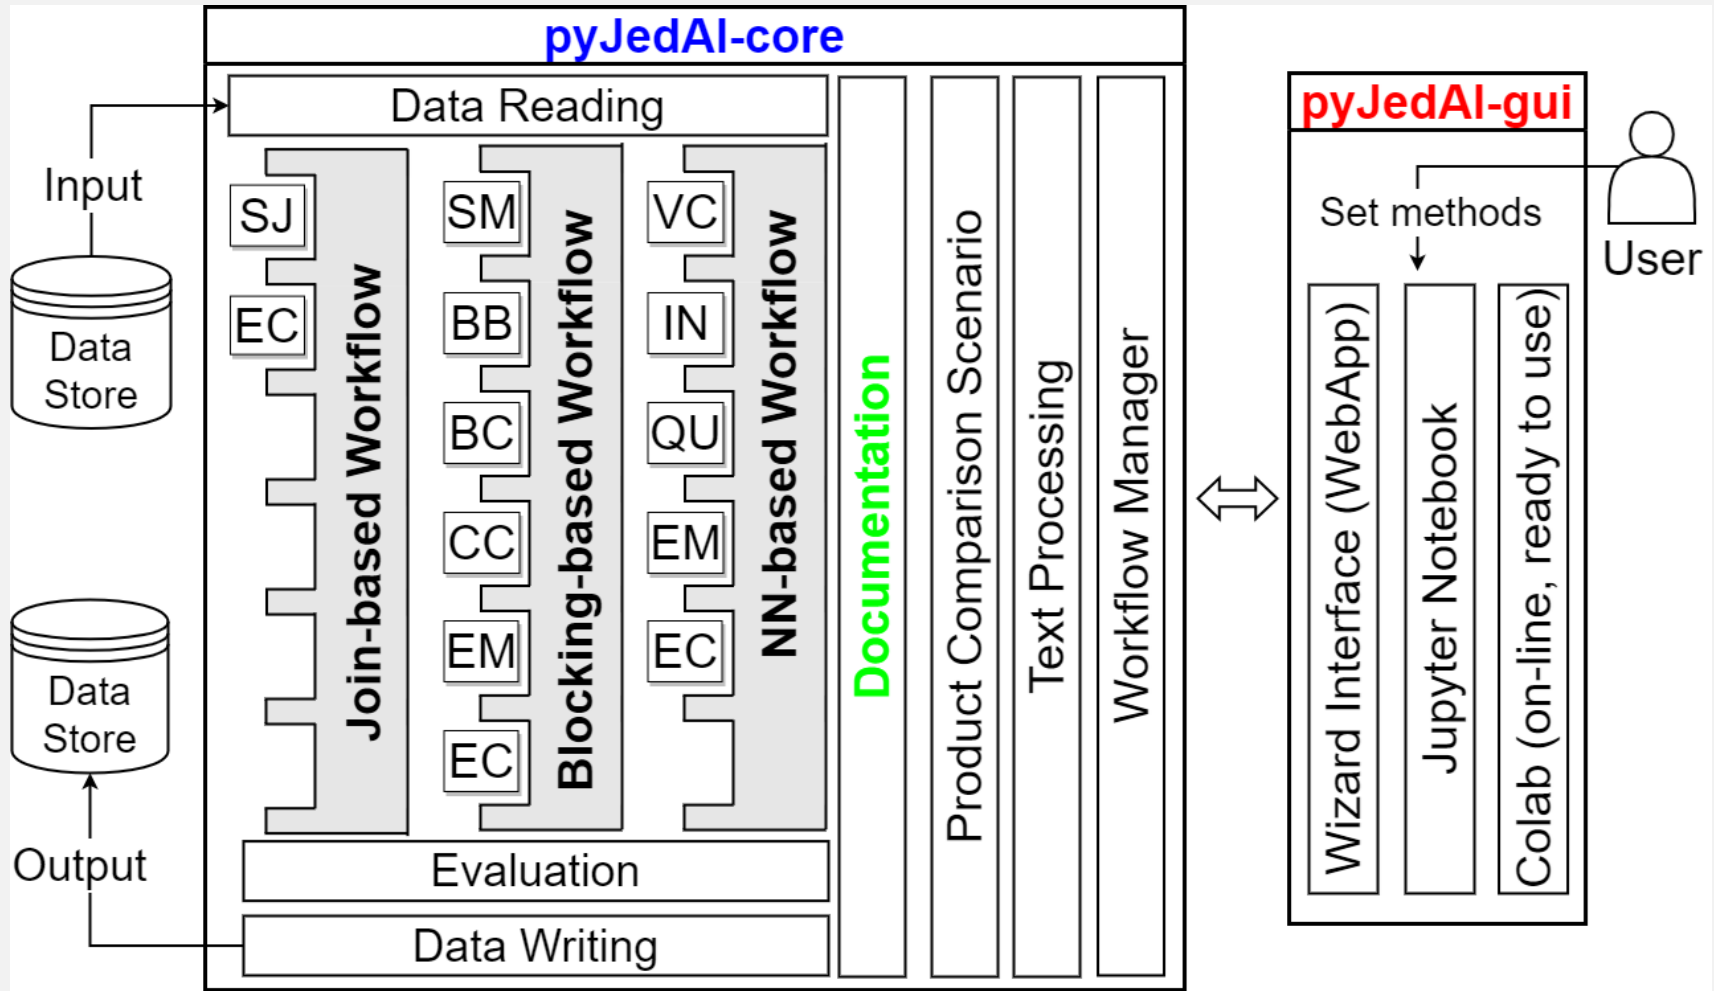
- Challenges and Final Remarks

# Our tool for ER – pyJedAI!

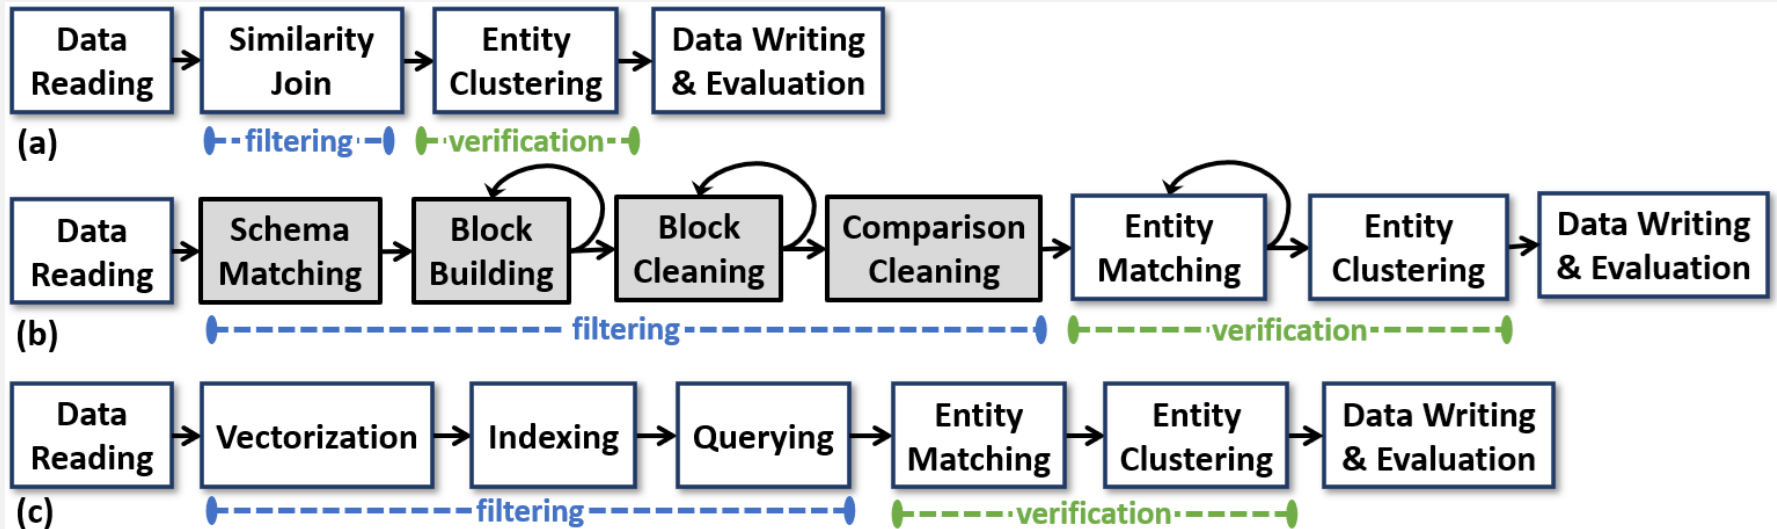
The logo for pyJedAI, featuring a stylized square icon with a smaller square inside, followed by the text "pyJedAI" in a bold, sans-serif font.

- A library of end-to-end ER workflows leveraging the Filtering-Verification framework
- *pyJedAI* is an open-source Python framework, supporting both experts and novice users, that leverages the latest breakthroughs in Deep Learning and NLP techniques, which are publicly available through the data science ecosystem
- Available at: <https://github.com/AI-team-UoA/pyJedAI>,
- Extends the **JedAI** tool that is implemented in Java  
(available at: <https://github.com/scify/JedAIToolkit>)

# pyJedAI Architecture



# 3 main Workflows

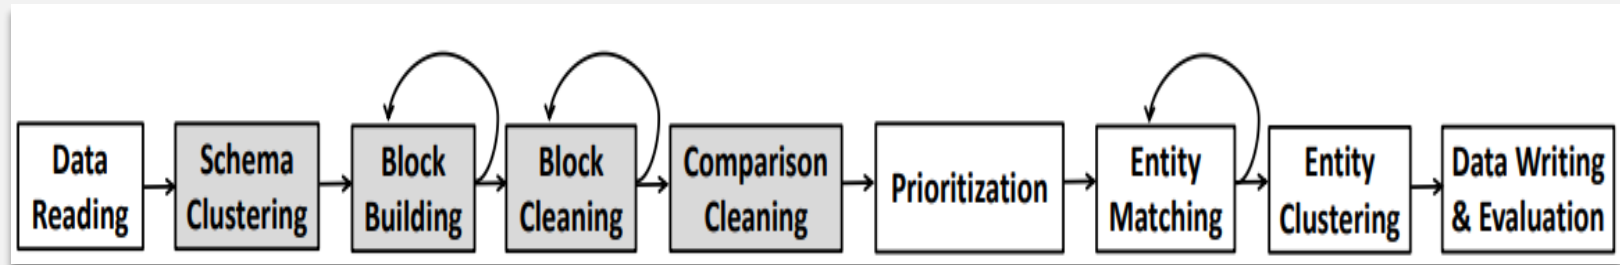


(a) Joins-based workflow

(b) Blocking-based workflow

(c) NN-based with embeddings workflow


# Blocking-based workflow



**Link to tutorial:**

<https://pyjedai.readthedocs.io/en/latest/tutorials/CleanCleanER.html>


# NN-based with embeddings workflow



**Link to tutorial:**

<https://pyjedai.readthedocs.io/en/latest/tutorials/pyTorchWorkflow.html>

# Joins-based workflow




**Link to tutorial:**

<https://pyjedai.readthedocs.io/en/latest/tutorials/SimilarityJoins.html>

# Install pyJedAI!

Scan QR and start  
entity-linking with  
**pyJedAI!**





- Introduction
- Generations: 1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup>, 5<sup>th</sup>
- Hands-on Session

## Part D: Challenges & Final Remarks

# Conclusions

---

Most promising works focus on:

## 1. Deep Learning

- Pros:
  - High accuracy
- Cons:
  - **High** training time
  - Too **many** training instances

## 2. Crowd-sourcing

- Pros:
  - High accuracy
- Cons:
  - High monetary **cost**
  - **Not scalable** to very large datasets

# Challenges

---

## Many challenges ahead

- Address shortcomings of Deep Learning
  - e.g., transfer learning for reducing labelling cost
- Cover gaps
  - e.g., incremental ER for semi-structured data
- New domains
  - e.g., adapt aforementioned techniques to **privacy-preserving** Entity Resolution

# ER Systems

---

- Literature focuses on **stand-alone** methods
- More emphasis on **end-to-end** systems
  - Examples: **Magellan** [1], **JedAI** [2]
  - Partially cover the 4 generations
  - More efforts meeting the following requirements [1,3]:
    - **open-source, extensible** systems
    - process data of **any structuredness**
    - **no coding!** for users
    - guidelines for creating effective solutions
    - covers the entire end-to-end pipeline exploit
    - a wide range of techniques

# Automatic Configuration

---

## **Facts:**

- Several parameters in every method
  - Applies to all generations and workflow steps
- Performance sensitive to internal configuration
- Manual fine-tuning required

## **Open Research Directions:**


- Plug-and-play methods
- Data-driven configuration

- Introduction
- Generations: 1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup>, 4<sup>th</sup>, 5<sup>th</sup>
- Hands-on Session
- Challenges &  
Final Remarks



# Thank you!

information & material  
related to the tutorial  
is available online



<https://edu.nl/97b8v>

# References

---

1. P. Konda, S. Das, P. S. G. C., A. Doan, A. Ardalan, J. R. Ballard, H. Li, F. Panahi, H. Zhang, J. F. Naughton, S. Prasad, G. Krishnan, R. Deep, and V. Raghavendra. Magellan: Toward building entity matching management systems. *PVLDB*, 9(12):1197–1208, 2016.
2. G. Papadakis, L. Tsekouras, E. Thanos, G. Giannakopoulos, T. Palpanas, and M. Koubarakis. Domain- and structureagnostic end-to-end entity resolution with jedai. *SIGMOD Record*, 48(4):31, 2019.
3. B. Golshan, A. Y. Halevy, G. A. Mihaila, and W. Tan. Data integration: After the teenage years. In *PODS*, pages 101–106, 2017.